


# Scaling BlockChain by Sharding



Debdoot Roy Chowdhury (23M0675)


Guided By: Prof. Vinay J. Ribeiro

# What is a Blockchain?

Blockchain is a decentralized, distributed, shared and immutable digital ledger that records data across many peers.


# What is a Blockchain?

Blockchain is a **decentralized**, **distributed**, **shared** and **immutable** digital ledger that records data across many peers.




# What is a Blockchain?

Blockchain is a **decentralized**, **distributed**, **shared** and **immutable** digital ledger that records data across many peers.




# What is a Blockchain?

Blockchain is a **decentralized, distributed, shared** and **immutable** digital **ledger** that records data across many peers.



# What is a Blockchain?


Blockchain is a **decentralized, distributed, shared** and **immutable** digital ledger that records data across many peers.





## What is Mining?

Users compete to create a block for getting rewards by solving a computationally intensive puzzle which takes significant computational effort.

# The Scalability Trilemma





# The Scalability Trilemma



Vitalik Buterin  
Founder of ETH




# The Scalability Trilemma




Vitalik Buterin  
Founder of ETH

Not possible  
to get all  
three


# The Scalability Trilemma




# The Scalability Trilemma



# The Scalability Trilemma



# The Scalability Trilemma



**NEED**

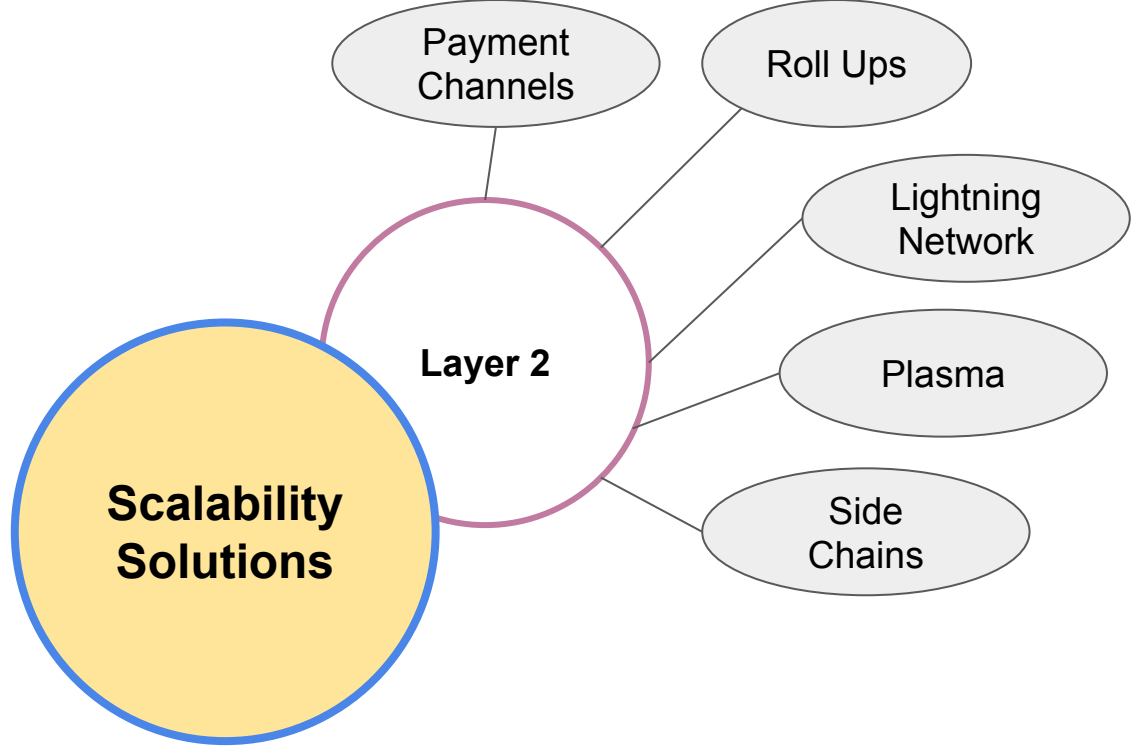
*More Speed*

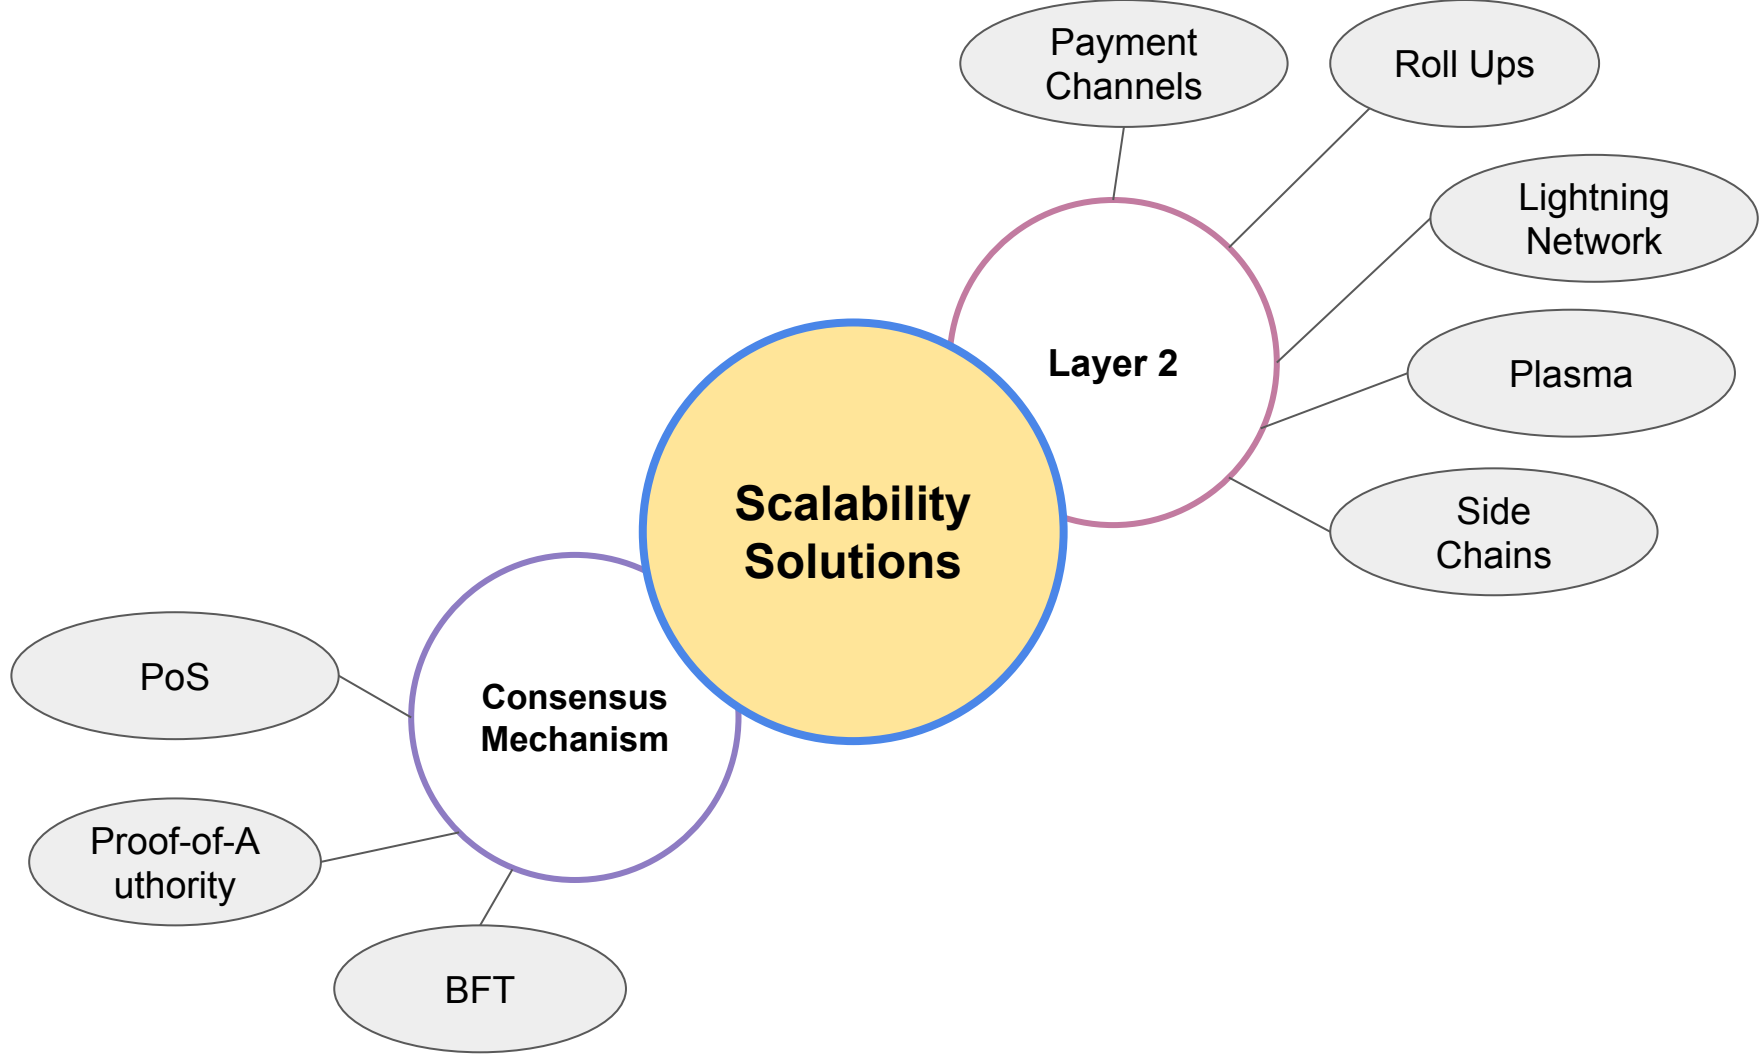
*Less Storage Overhead*

*More Miners = More Speed*

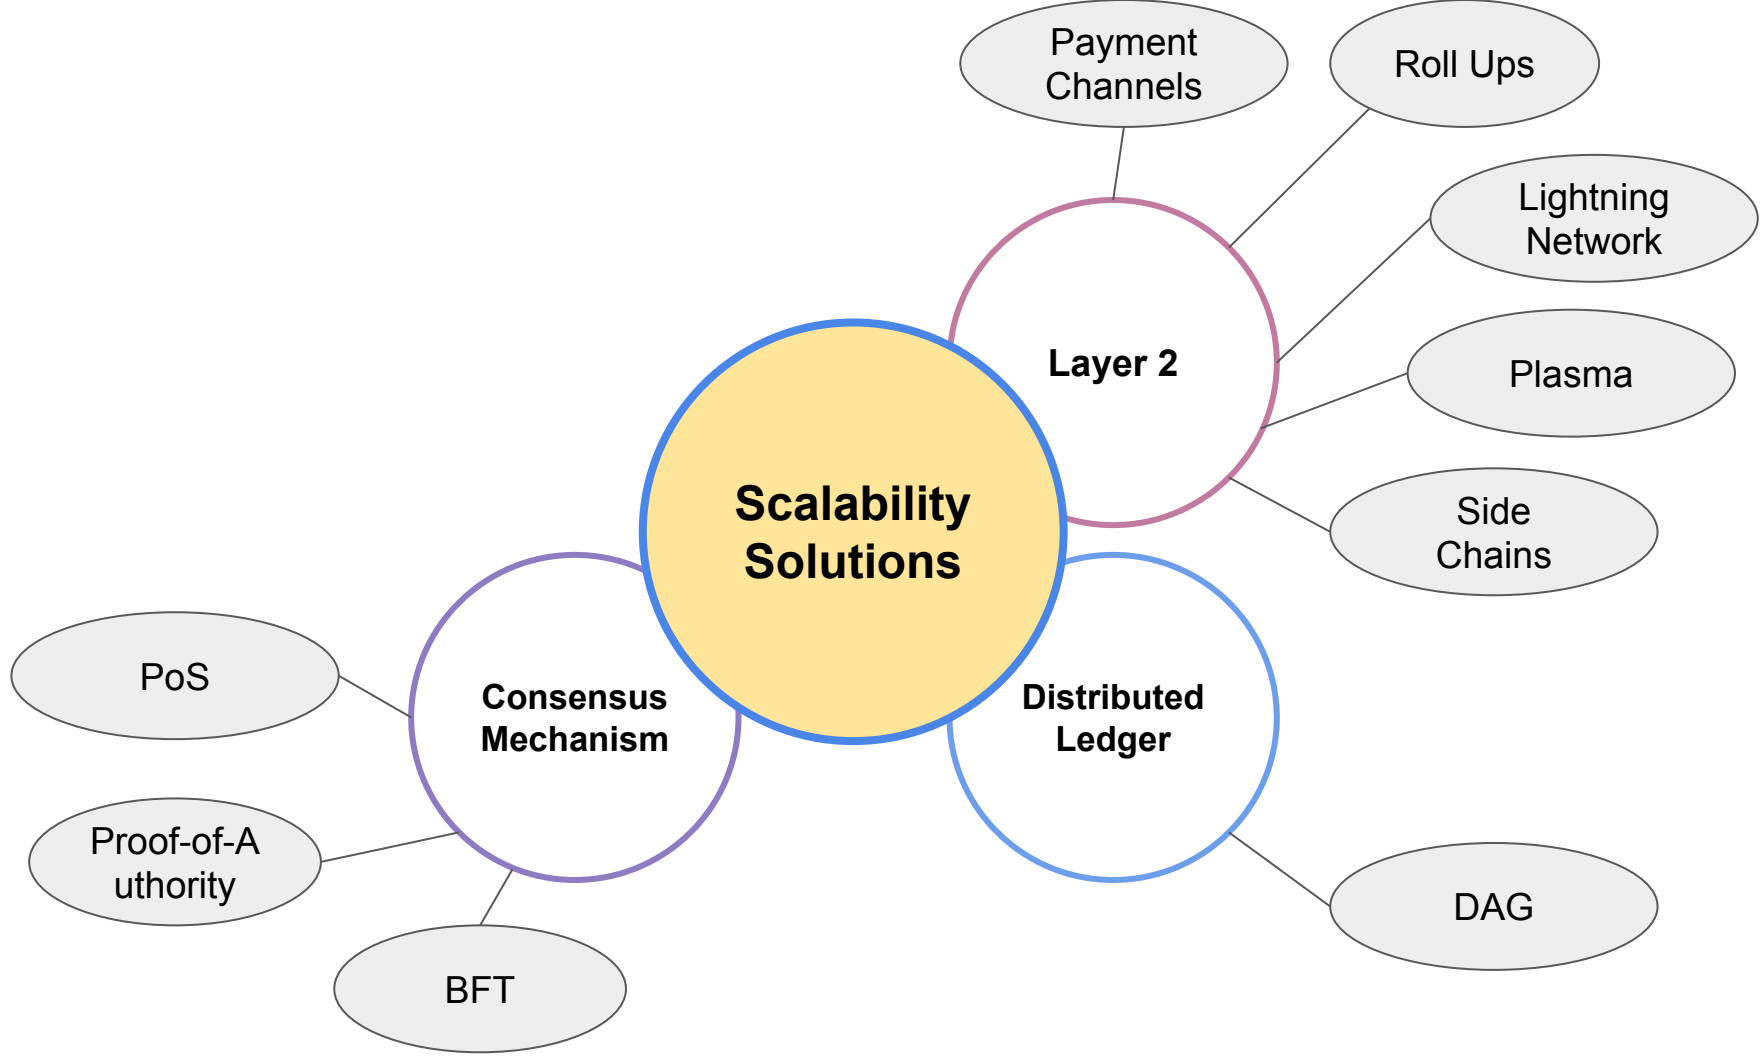


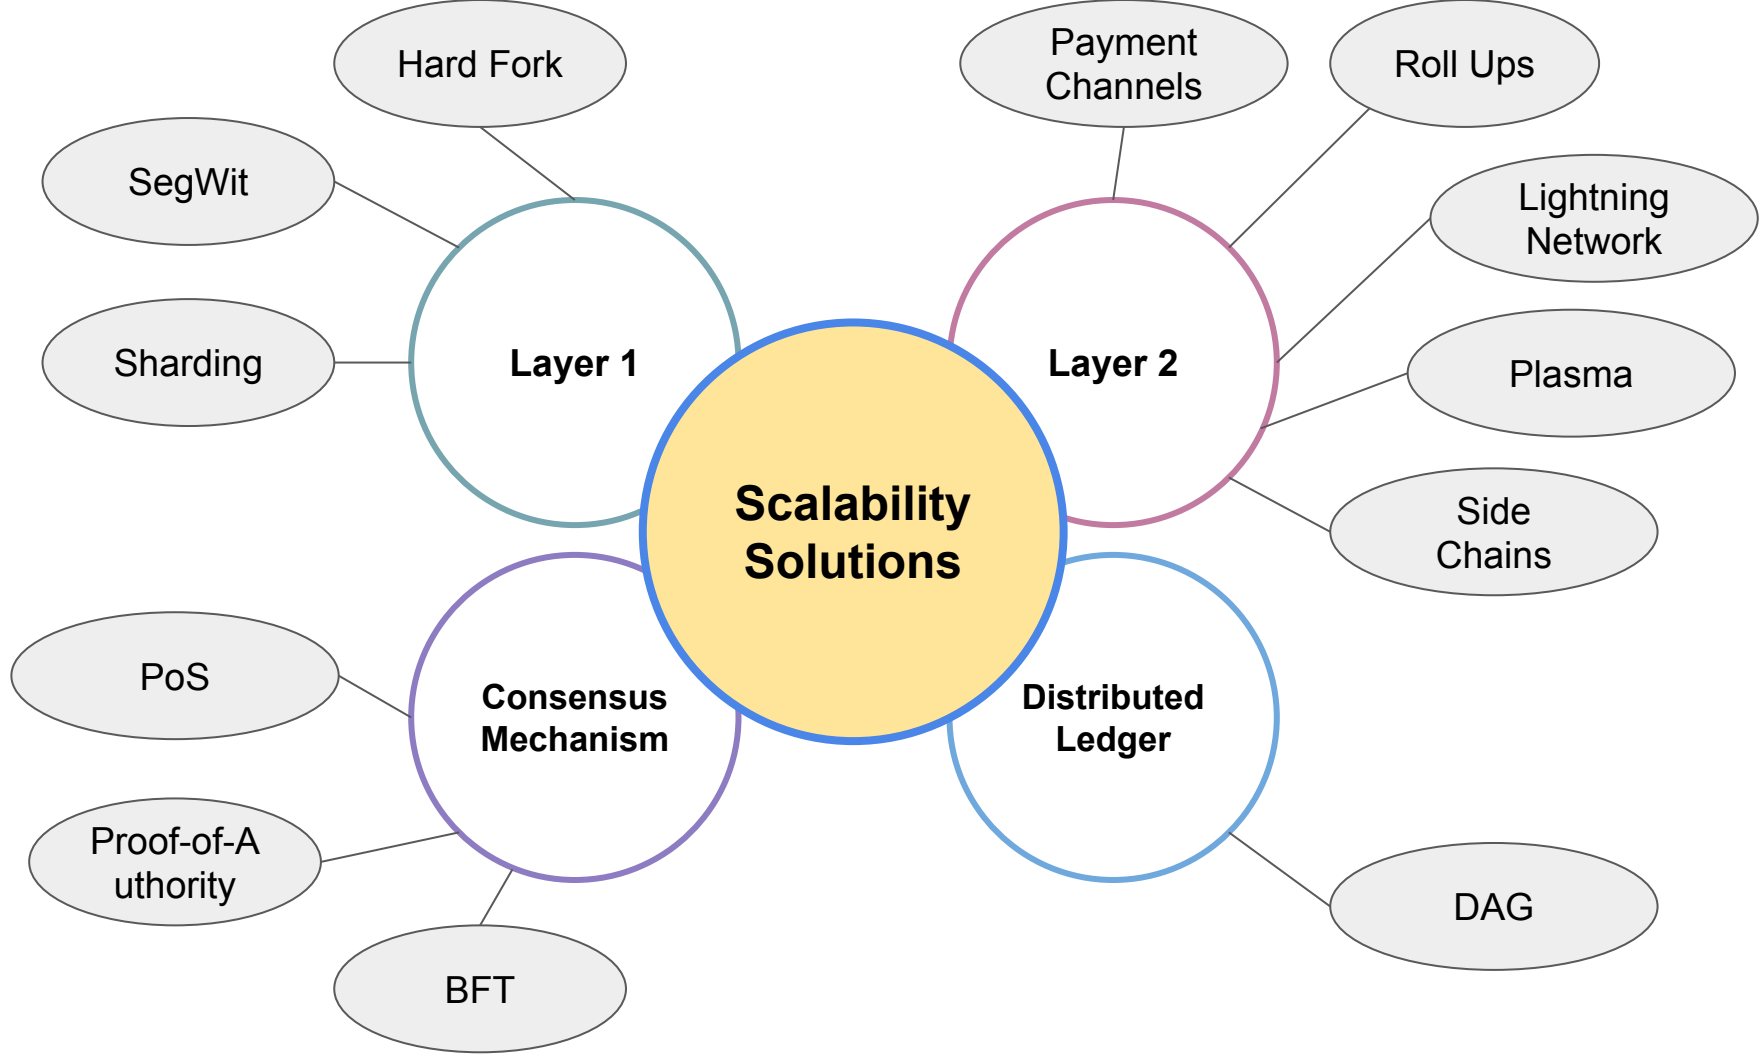
**Scalability  
Solutions**

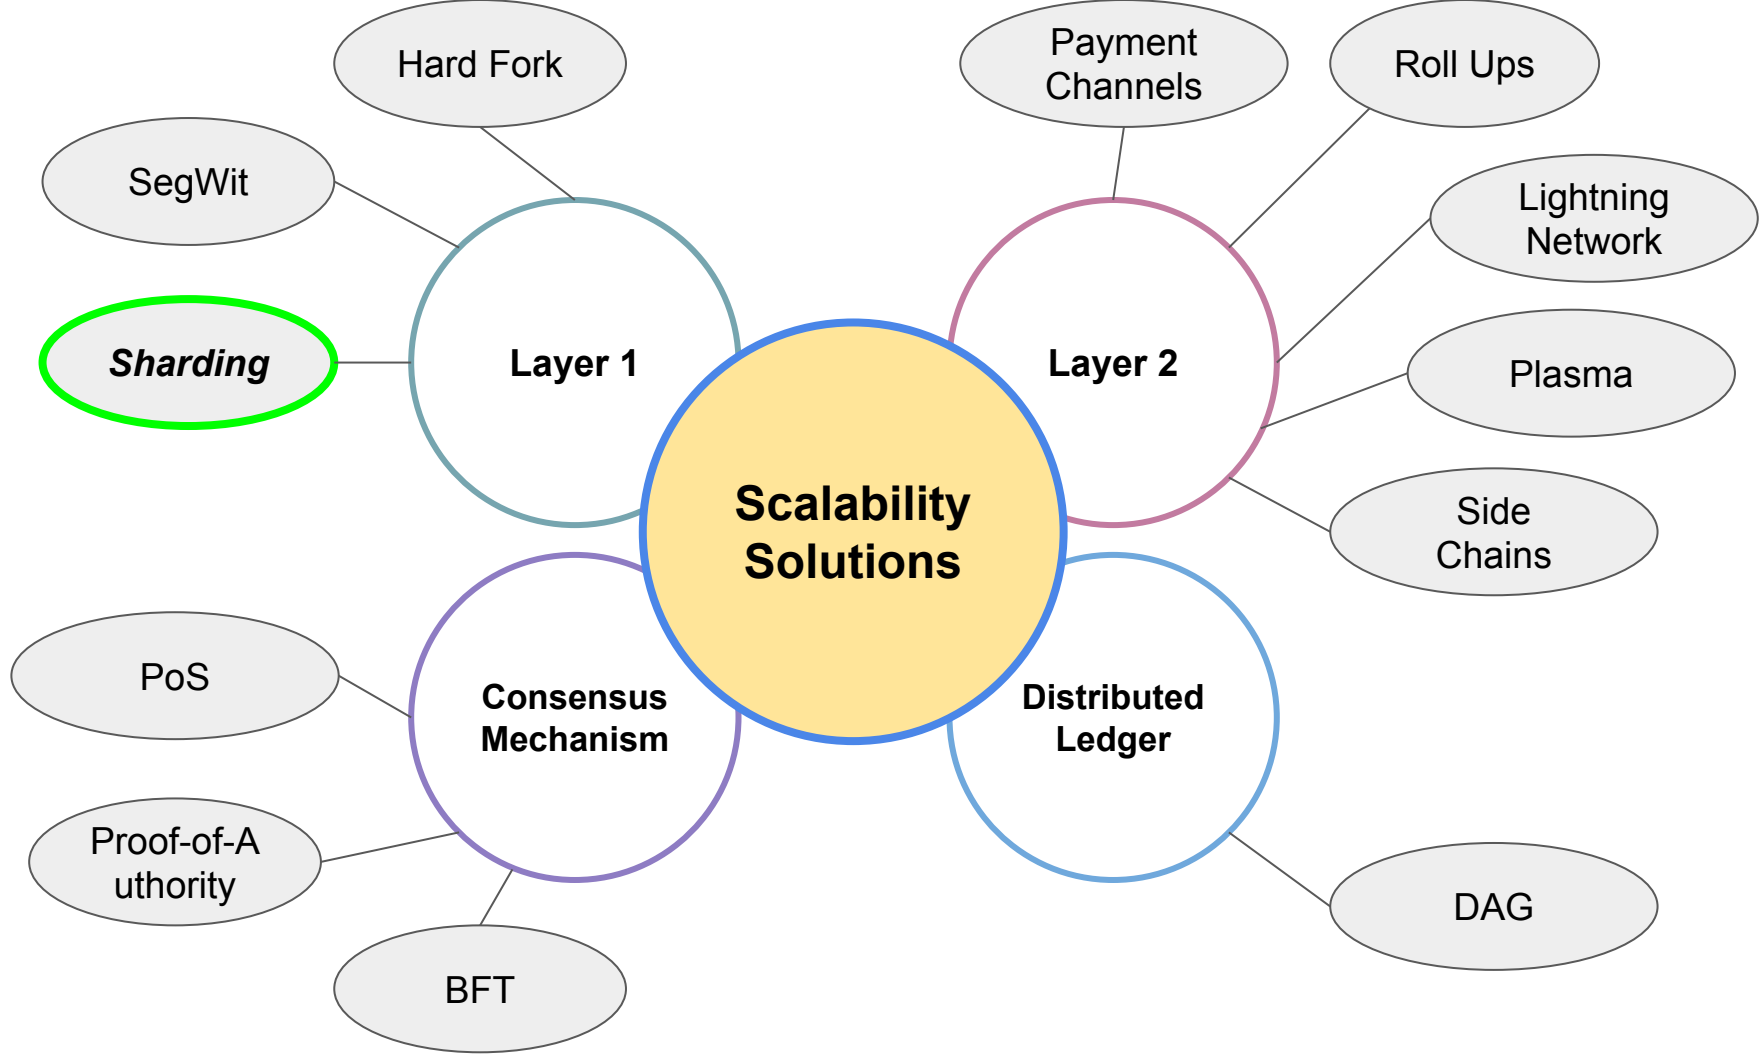













# What is Sharding?

Partitions the network into smaller groups called **shards**, allowing *parallel processing* of transactions to improve overall throughput.


# What is Sharding?

Partitions the network into smaller groups called **shards**, allowing *parallel processing* of transactions to improve overall throughput.



# What is Sharding?

Partitions the network into smaller groups called **shards**, allowing *parallel processing* of transactions to improve overall throughput.

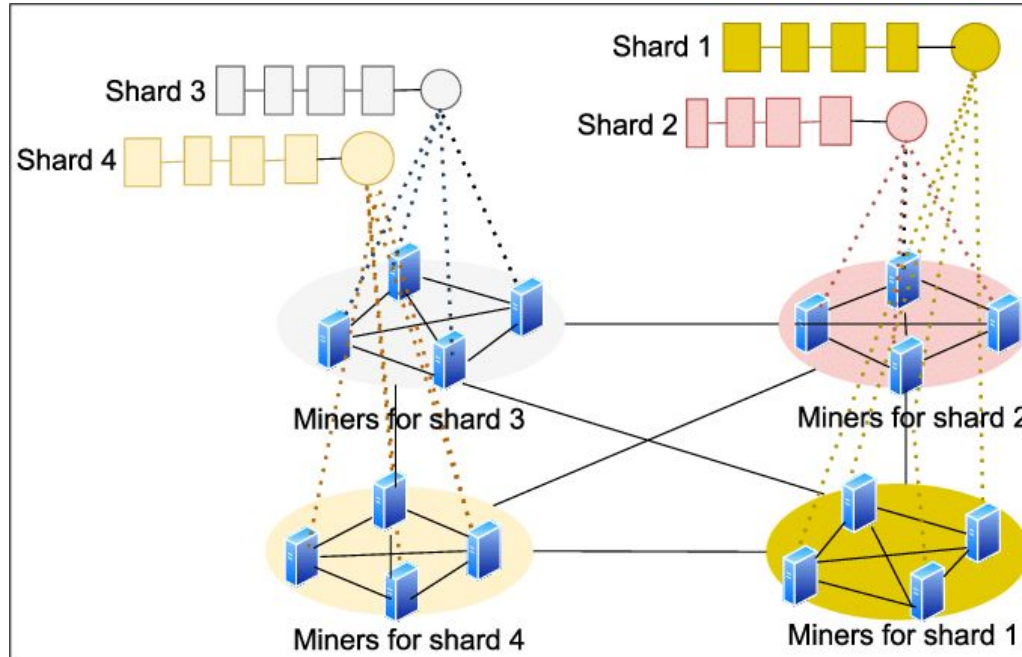


Does 4 parallel chains means 4x speed?



# What is Sharding?


Partitions the network into smaller groups called **shards**, allowing *parallel processing* of transactions to improve overall throughput.



Does 4 parallel chains means 4x speed?




Depends






Intra-Shard  
Communication

Miners or  
Validators




Intra-Shard  
Communication

Miners or  
Validators




Intra-Shard  
Communication

Miners or  
Validators




Cross-Shard  
Communication


# Challenges in Sharding




# Challenges in Sharding




# Challenges in Sharding



# Challenges in Sharding




# Challenges in Sharding

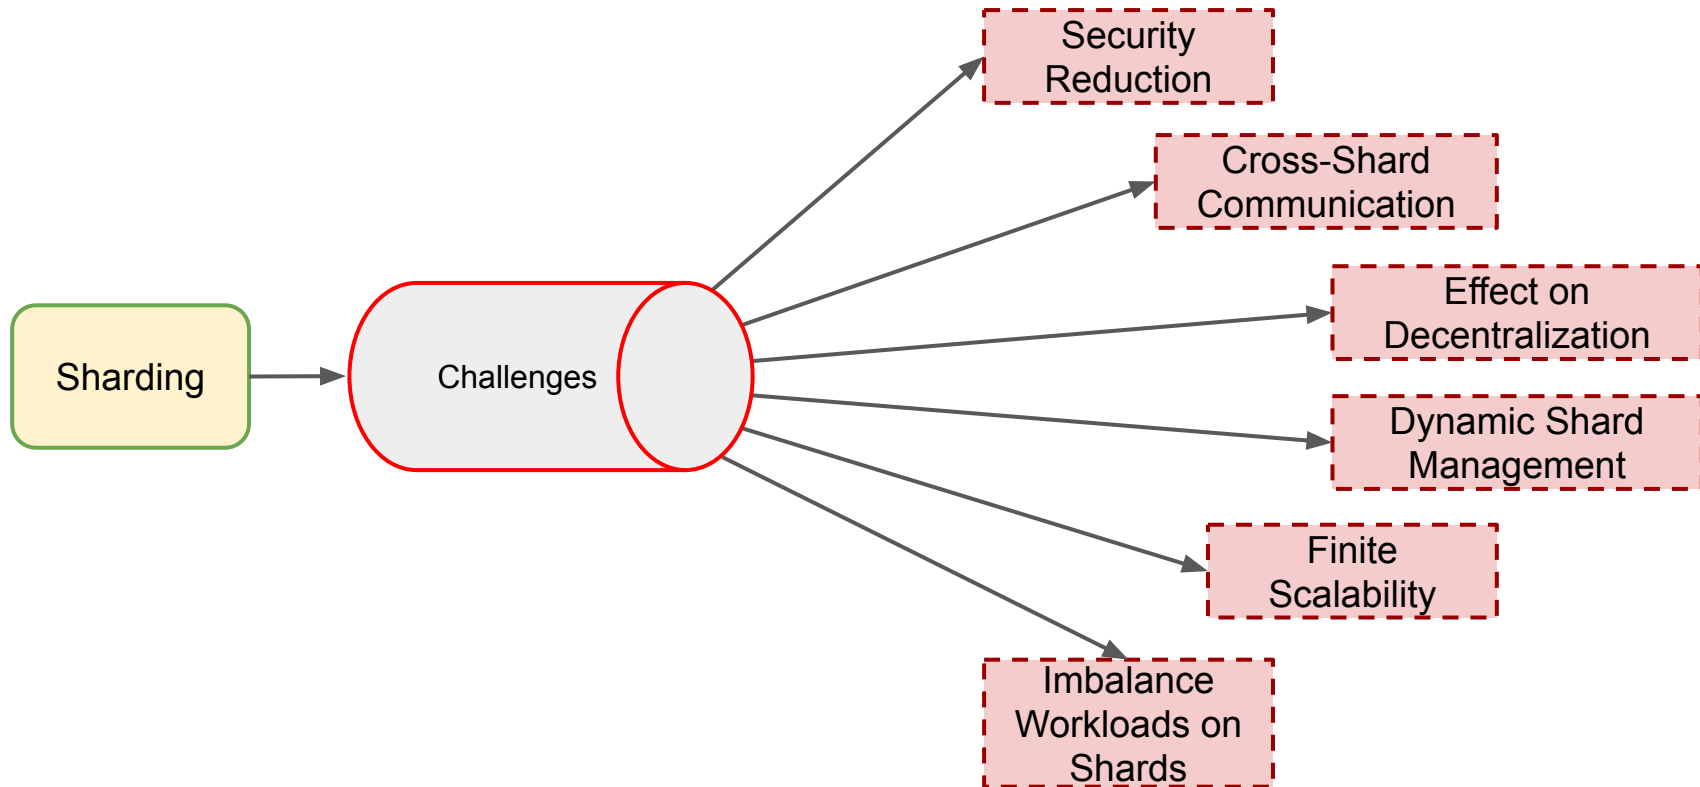




# Challenges in Sharding



# Challenges in Sharding



# Security Reduction


In a protocol where majority votes are considered as final decision, over 50% of corrupted nodes will always yield a corrupted block.

# Security Reduction

In a protocol where majority votes are considered as final decision, over 50% of corrupted nodes will always yield a corrupted block.

But, getting over 50% of the entire network is HARD!

X validators building one chain.  
**Need to corrupt 0.51x**




# Security Reduction

In a protocol where majority votes are considered as final decision, over 50% of corrupted nodes will always yield a corrupted block.

But, getting over 50% of the entire network is HARD!

**Famously known  
as 51% attack.**


X validators building one chain.  
**Need to corrupt 0.51x**



# Security Reduction

But, getting over 50% of a single shard is EASY!

X validators building 10 chains  
Need to corrupt 0.051x




# Security Reduction

But, getting over 50% of a single shard is EASY!


**Famously known  
as 1% attack.**

X validators building 10 chains  
Need to corrupt 0.051x




# Techniques of Sharding

## 1) Network Sharding



Shards




Nodes

Nodes here  
are the Miners  
or Validators  
of the Network




# Techniques of Sharding

## 1) Network Sharding



Shards




Nodes here  
are the Miners  
or Validators  
of the Network

Nodes

# Techniques of Sharding

Nodes here  
are the Miners  
or Validators  
of the Network

## 1) Network Sharding




Shards


Nodes

# Techniques of Sharding

## 2) Transaction Sharding




Shards



Transactions

# Techniques of Sharding

## 2) Transaction Sharding



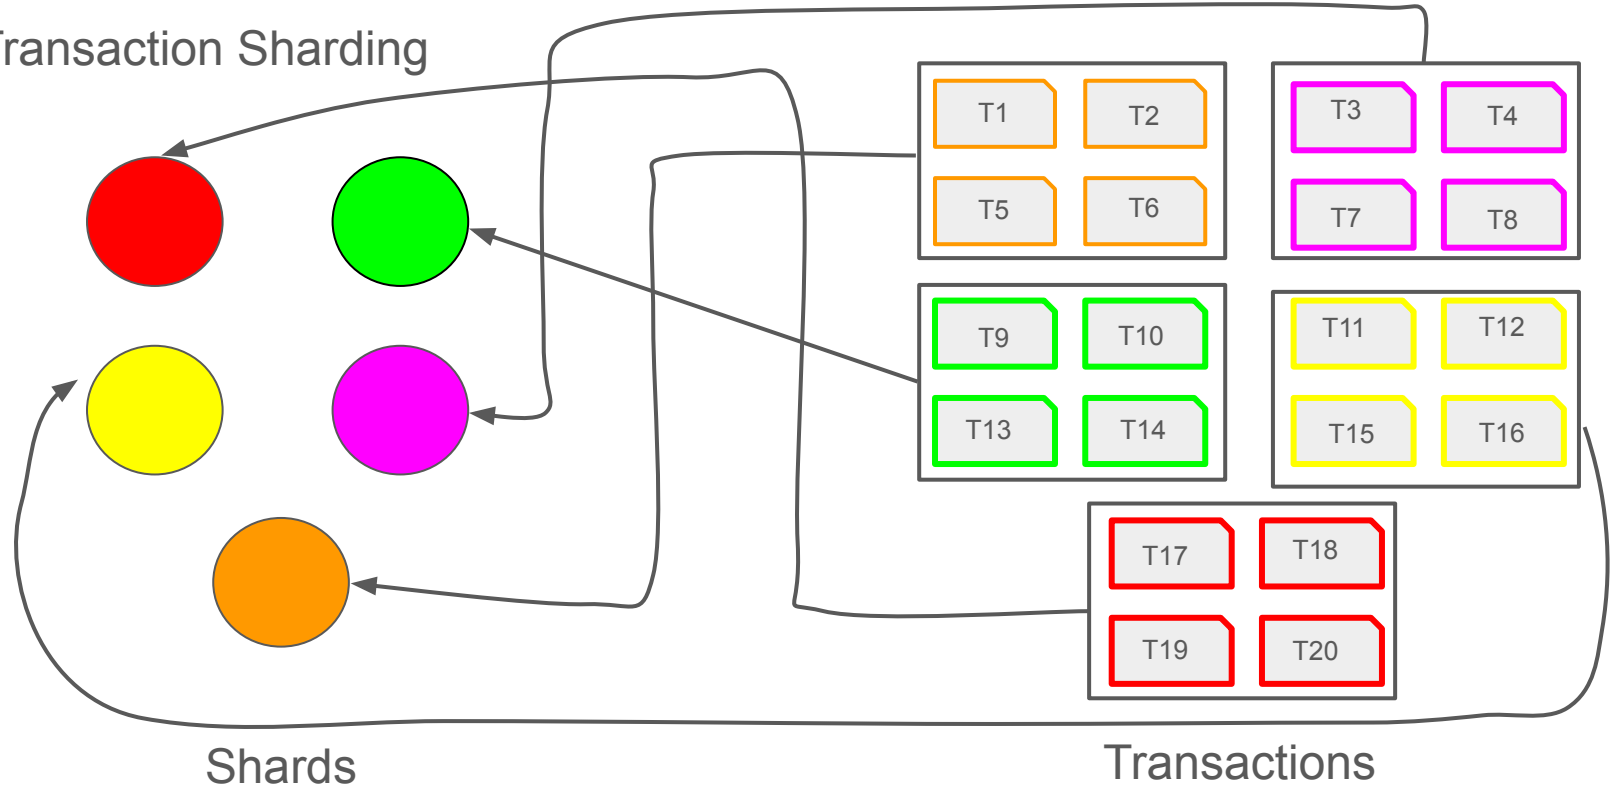
Shards



Transactions

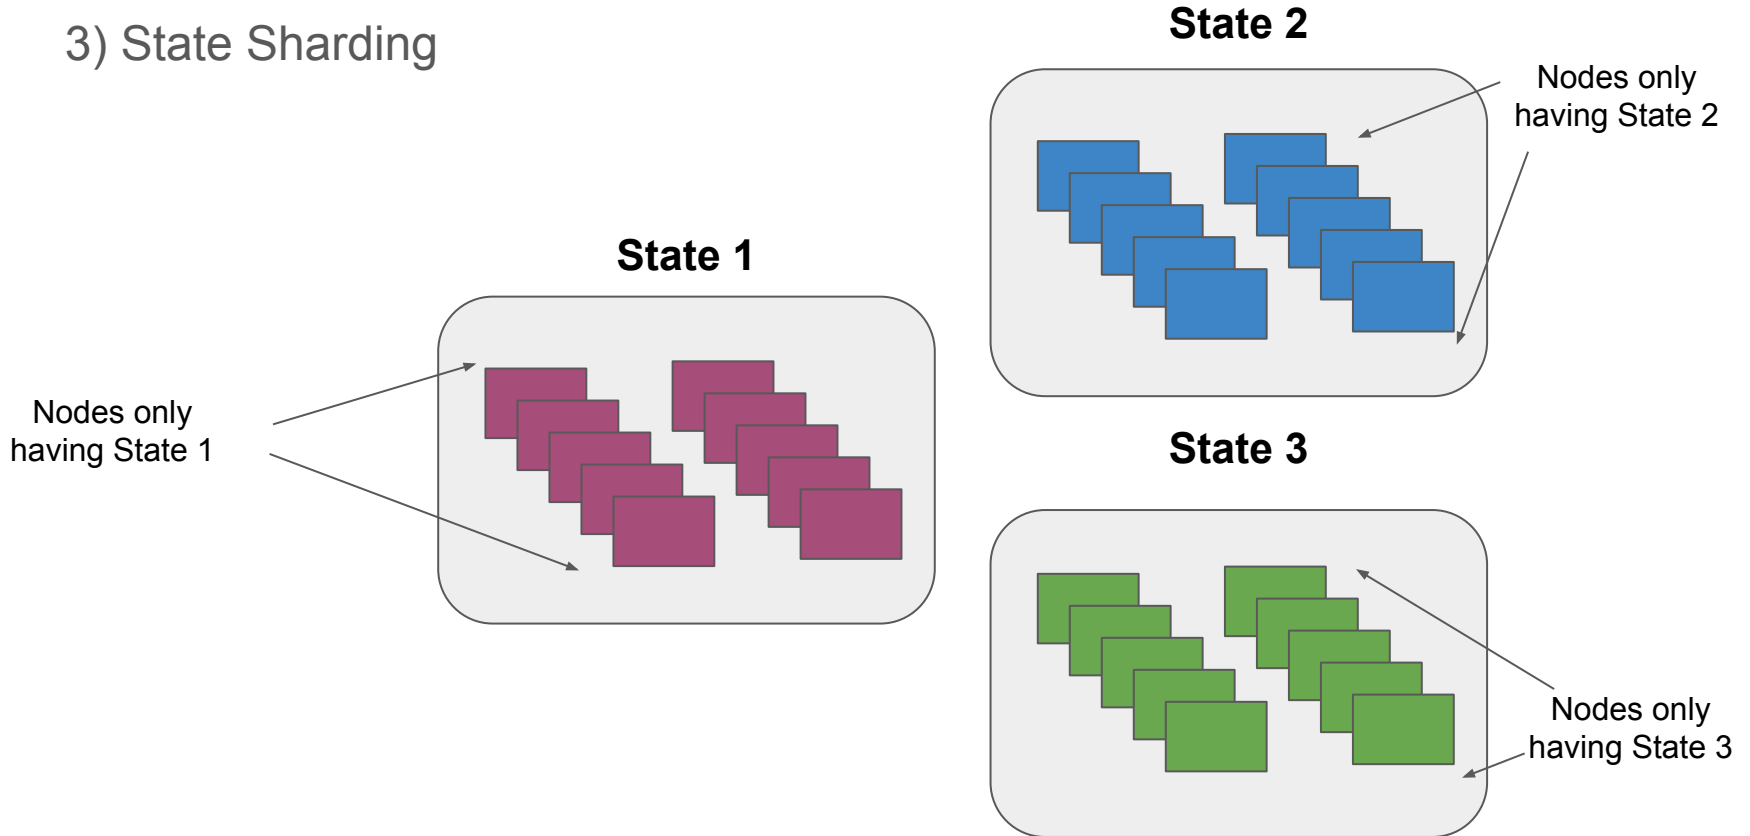
# Techniques of Sharding

## 2) Transaction Sharding



# Techniques of Sharding

## 3) State Sharding



# Metrics for evaluating a Sharding Solution

Scalability

Security

Fault  
Tolerance

Latency

Load  
Balancing

Throughput

Adaptability  
and  
Flexibility

Atomicity and Overhead of  
Cross-Shard Transactions

# Metrics for evaluating a Sharding Solution

Scalability

Security

Fault  
Tolerance

Latency

Load  
Balancing

Throughput

Adaptability  
and  
Flexibility

Atomicity and Overhead of  
Cross-Shard Transactions



# Steps in Sharding

1) Identity Establishment

# Steps in Sharding

## 1) Identity Establishment

Nodes need to establish their identity using some method like PoX.

# Steps in Sharding

## 1) Identity Establishment

Nodes need to establish their identity using some method like PoX.

Can i make a sybil attack?



# Steps in Sharding


## 1) Identity Establishment

Nodes need to establish their identity using some method like PoX.

Can i make a sybil attack?



Oh No! They are using PoW to establish identity.



# Steps in Sharding

## 1) Identity Establishment

Nodes need to establish their identity using some method like PoX.

Can i make a sybil attack?



Oh No! They are using PoW to establish identity.



## 2) Shard Assignment and Setup

# Steps in Sharding

## 1) Identity Establishment

Nodes need to establish their identity using some method like PoX.

Can i make a sybil attack?



Oh No! They are using PoW to establish identity.



## 2) Shard Assignment and Setup

Nodes and transaction randomly distributed among shard based on their identity.

# Steps in Sharding


## 1) Identity Establishment

Nodes need to establish their identity using some method like PoX.

Can i make a sybil attack?




Oh No! They are using PoW to establish identity.



## 2) Shard Assignment and Setup

Nodes and transaction randomly distributed among shard based on their identity.

All Nodes



# Steps in Sharding

## 1) Identity Establishment

Nodes need to establish their identity using some method like PoX.

Can i make a sybil attack?

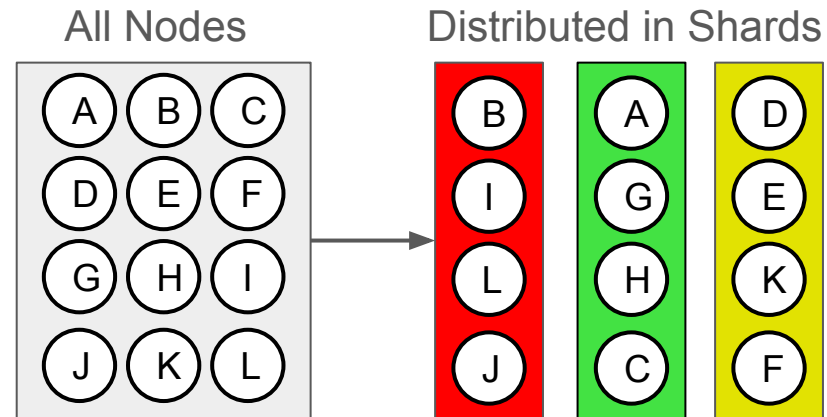


Oh No! They are using PoW to establish identity.



## 2) Shard Assignment and Setup

Nodes and transaction randomly distributed (non-deterministically) among shard based on their identity.





# Steps in Sharding

3) Consensus Determination

# Steps in Sharding

## 3) Consensus Determination

Consensus needs to be reached for both intra-shard and cross-shard communication.

# Steps in Sharding

## 3) Consensus Determination

Consensus needs to be reached for both intra-shard and cross-shard communication.

S1



# Steps in Sharding

## 3) Consensus Determination

Consensus needs to be reached for both intra-shard and cross-shard communication.


S1



# Steps in Sharding

## 3) Consensus Determination


Consensus needs to be reached for both intra-shard and cross-shard communication.



# Steps in Sharding

## 3) Consensus Determination

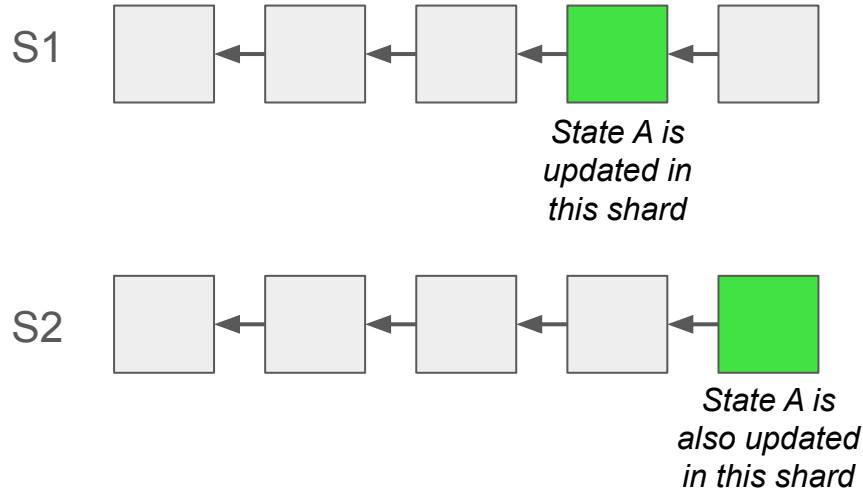
Consensus needs to be reached for both intra-shard and cross-shard communication.



# Steps in Sharding

## 3) Consensus Determination

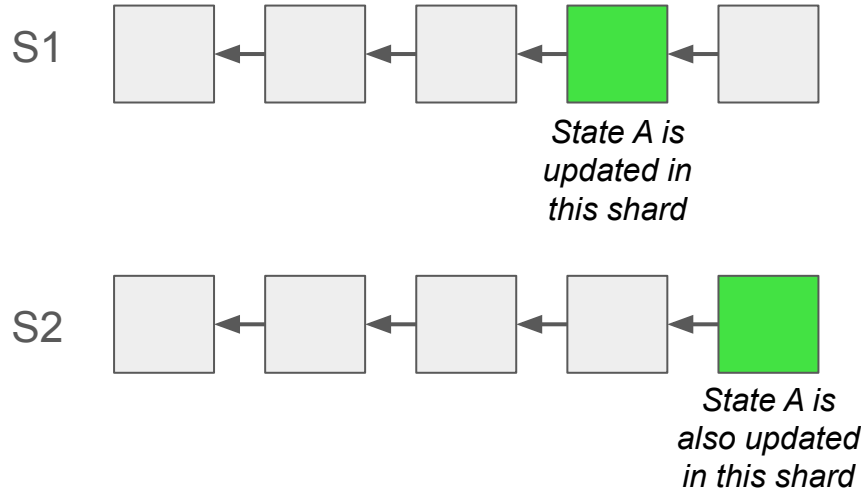
Consensus needs to be reached for both intra-shard and cross-shard communication.



# Steps in Sharding

## 3) Consensus Determination

Consensus needs to be reached for both intra-shard and cross-shard communication.




## 4) Shard Reconfiguration



# Steps in Sharding

## 3) Consensus Determination

Consensus needs to be reached for both intra-shard and cross-shard communication.




## 4) Shard Reconfiguration

Reconfiguration with redistribution of nodes among shard after each epoch to bolster security.

# Steps in Sharding


## 3) Consensus Determination

Consensus needs to be reached for both intra-shard and cross-shard communication.



## 4) Shard Reconfiguration


Reconfiguration with redistribution of nodes among shard after each epoch to bolster security.



# Steps in Sharding


## 3) Consensus Determination

Consensus needs to be reached for both intra-shard and cross-shard communication.



## 4) Shard Reconfiguration

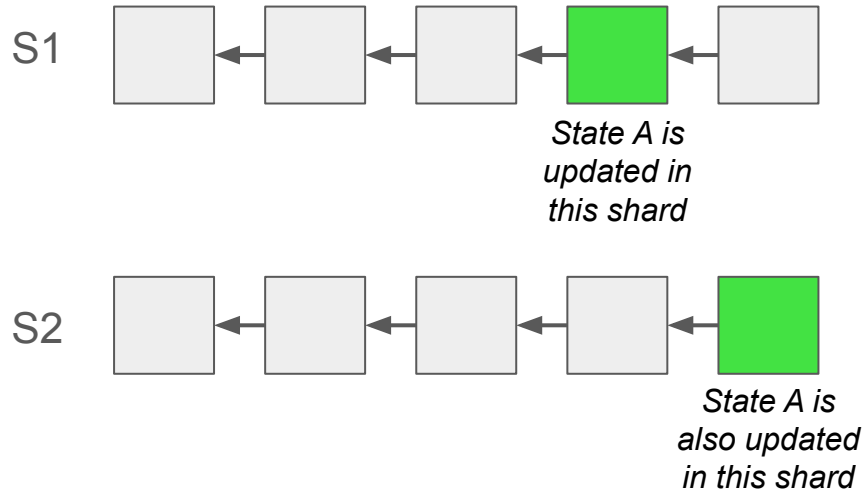
Reconfiguration with redistribution of nodes among shard after each epoch to bolster security.



# Steps in Sharding


## 3) Consensus Determination

Consensus needs to be reached for both intra-shard and cross-shard communication.



## 4) Shard Reconfiguration

Reconfiguration with redistribution of nodes among shard after each epoch to bolster security.



# Elastico [CCS'16]

First secure sharding protocol for open blockchain.



# Elastico [CCS'16]

First secure sharding protocol for open blockchain.




What problem is  
addressed?

# Elastico [CCS'16]

First secure sharding protocol for open blockchain.



What problem is addressed?



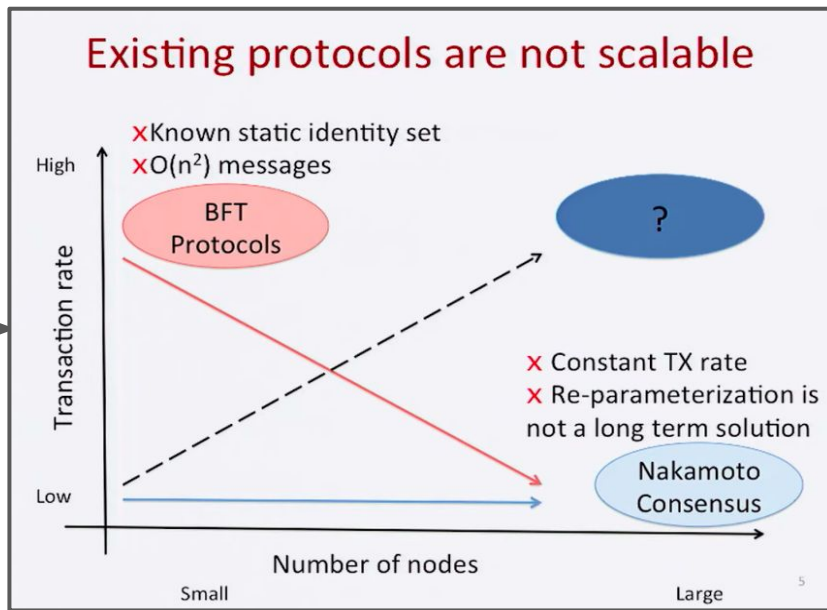
Luu, Loi, et al. "A secure sharding protocol for open blockchains."

# Elastico [CCS'16]

First secure sharding protocol for open blockchain.



What problem is addressed?



Can we make "n" smaller, so BFT protocols can run faster?

Luu, Loi, et al. "A secure sharding protocol for open blockchains."

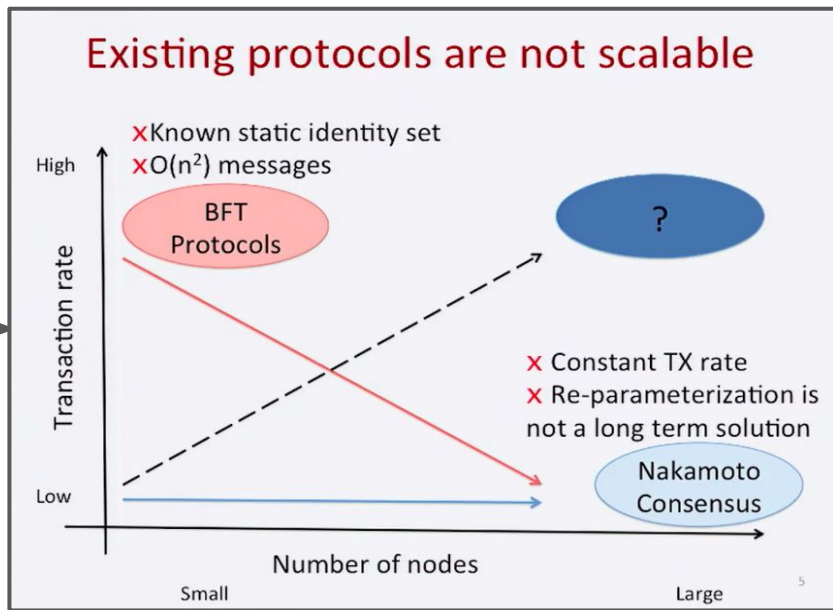


# Elastico [CCS'16]

First secure sharding protocol for open blockchain.



What problem is addressed?



Can we make "n" smaller, so BFT protocols can run faster?

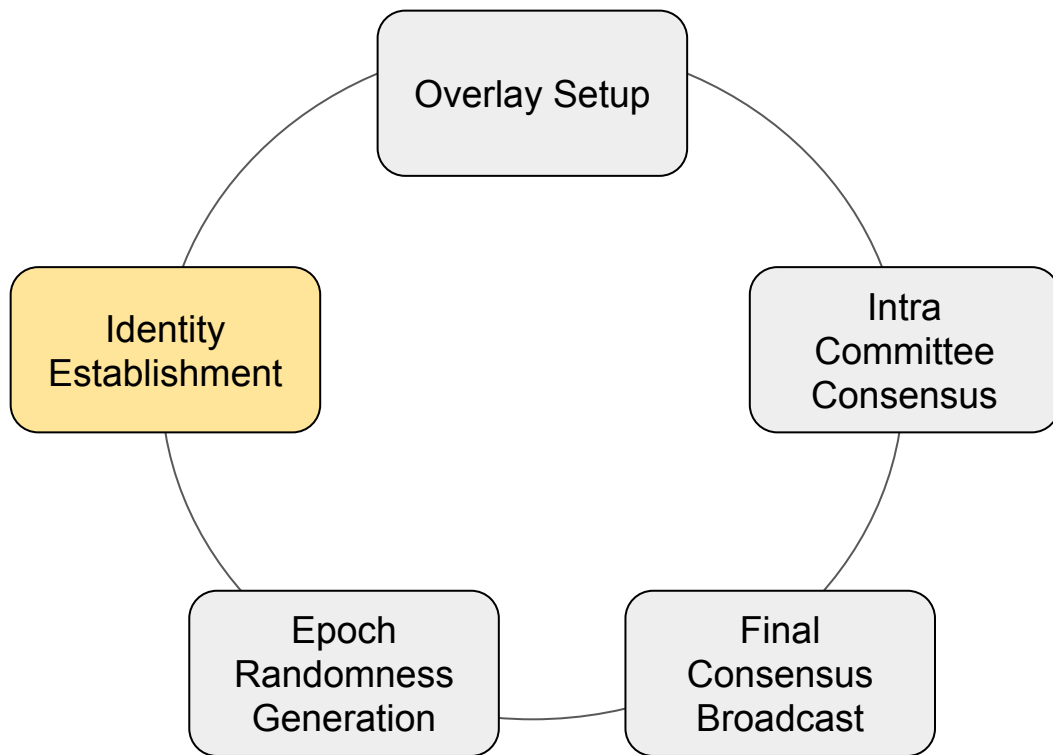


Sharding

Luu, Loi, et al. "A secure sharding protocol for open blockchains."

# Elastico

The protocol operates in Epochs.



# Elastico

## *Identity Establishment*

Nodes perform PoW to establish their identity.

# Elastico

## *Identity Establishment*

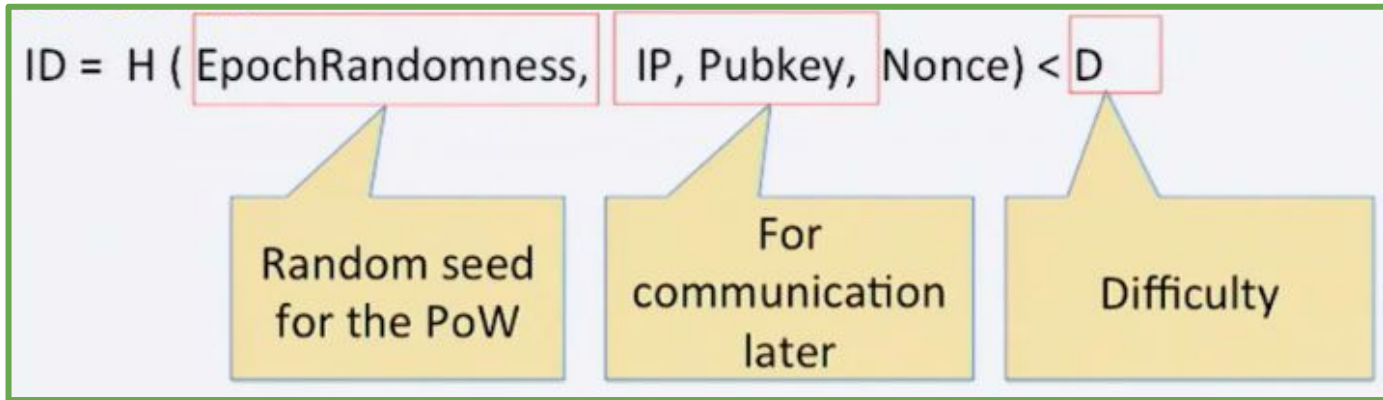
Nodes perform PoW to establish their identity.

$$\text{ID} = \text{H}(\text{EpochRandomness}, \text{IP, Pubkey, Nonce}) < \text{D}$$

# Elastico

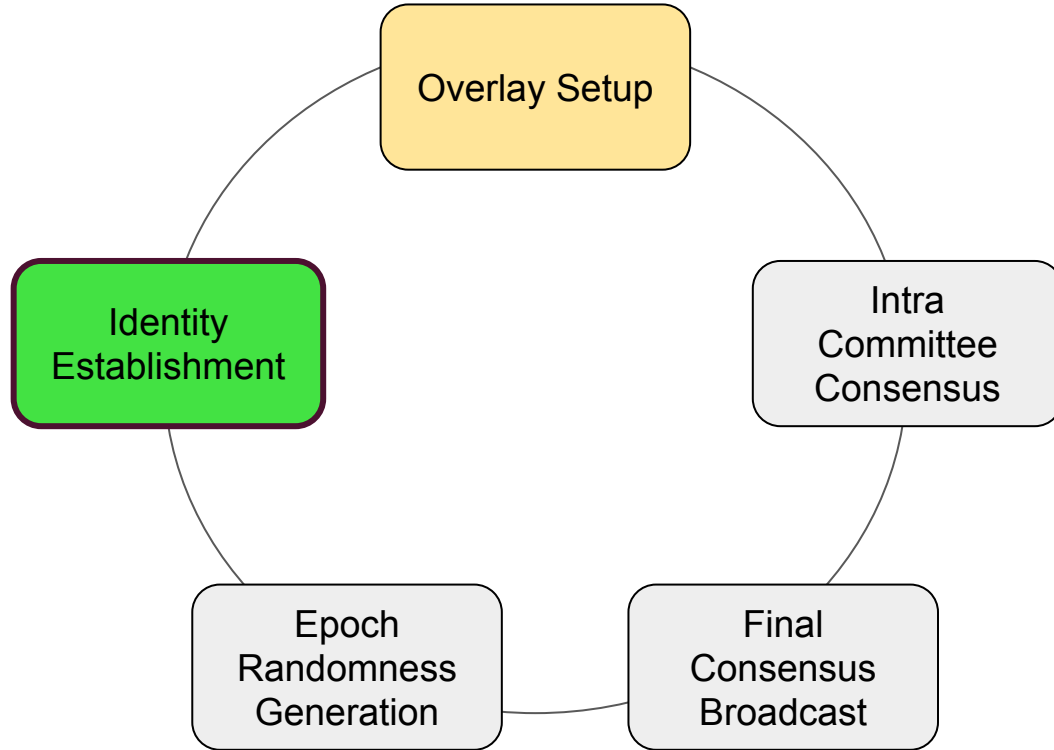
## *Identity Establishment*

Nodes perform PoW to establish their identity.



Luu, Loi, et al. "A secure sharding protocol for open blockchains."

# Elastico



# Elastico

## *Overlay Setup*

- Nodes are **assigned** to shard **according to** their **ID**.

# Elastico

## *Overlay Setup*

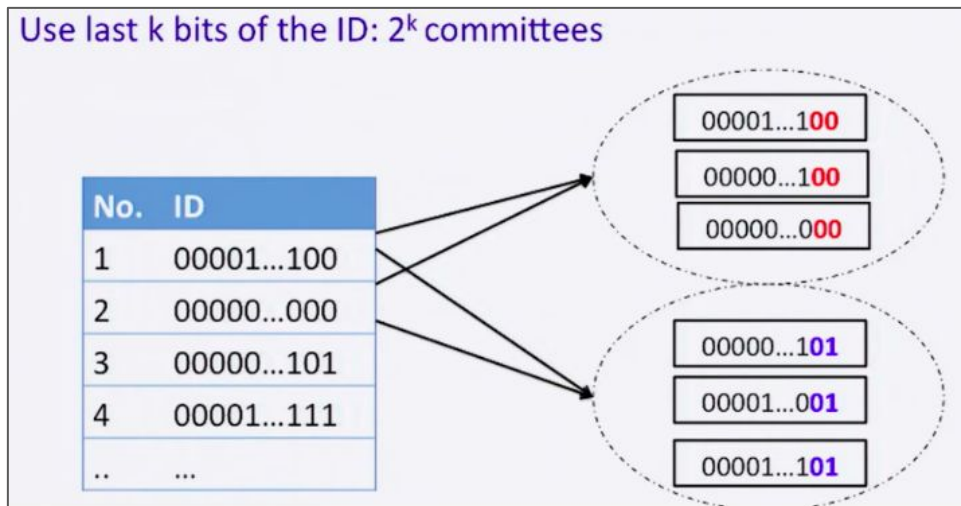
- Nodes are **assigned** to shard **according to** their **ID**.
- Elastico equitably allocates identities to  **$2^k$  committees**, utilizing the **last k bits of their identities (ID)**.



# Elastico

## Overlay Setup

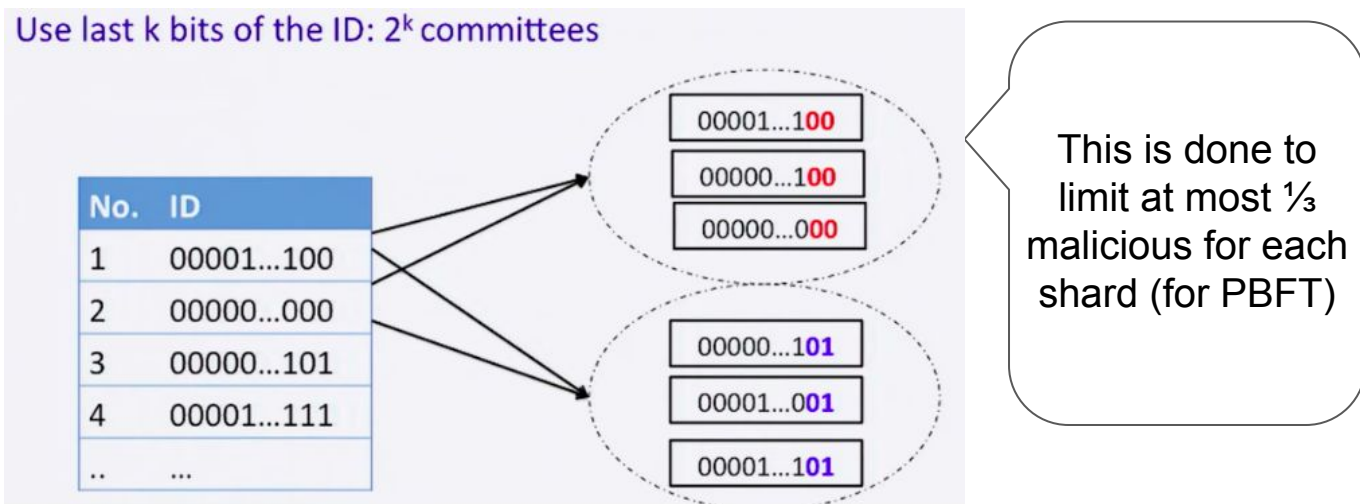
- Nodes are **assigned** to shard **according to** their **ID**.
- Elastico equitably allocates identities to  $2^k$  **committees**, utilizing the **last k bits** of their identities (ID).



# Elastico

## Overlay Setup

- Nodes are **assigned** to shard **according to** their **ID**.
- Elastico equitably allocates identities to  $2^k$  **committees**, utilizing the **last k bits** of their identities (ID).



How to get  
identities from  
other nodes?

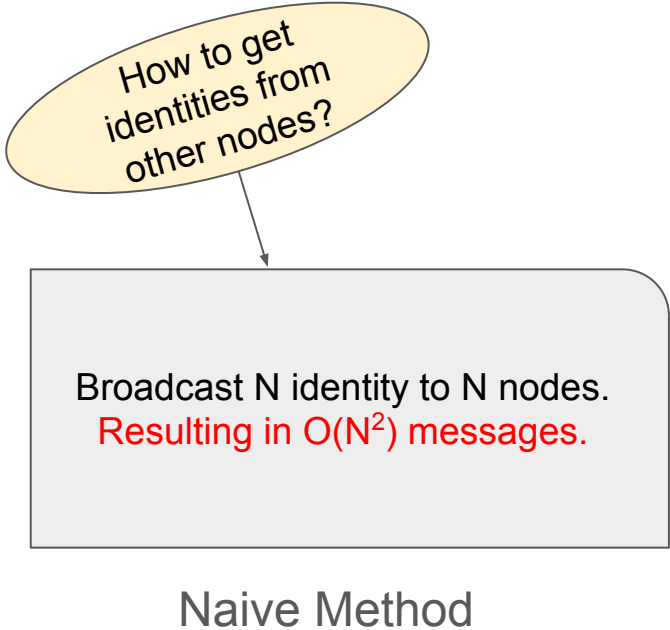
# Elastico

## *Overlay Setup*

# Elastico

## *Overlay Setup*

How to get  
identities from  
other nodes?



Broadcast N identity to N nodes.  
Resulting in  $O(N^2)$  messages.

Naive Method

# Elastico

## Overlay Setup

How to get  
identities from  
other nodes?

Broadcast N identity to N nodes.  
Resulting in  $O(N^2)$  messages.

Naive Method

1. First C identities becomes Directory Servers.

# Elastico

## Overlay Setup


How to get  
identities from  
other nodes?

Broadcast N identity to N nodes.  
Resulting in  $O(N^2)$  messages.

Naive Method

1. First C identities becomes Directory Servers.
2. Each node after establishing identity sends identity to nearest Directory Server.

N Nodes



# Elastico

## Overlay Setup


How to get  
identities from  
other nodes?

Broadcast N identity to N nodes.  
Resulting in  $O(N^2)$  messages.

Naive Method

1. First C identities becomes Directory Servers.
2. Each node after establishing identity sends identity to nearest Directory Server.

N Nodes



Directory Servers

# Elastico


## Overlay Setup

How to get identities from other nodes?

Broadcast N identity to N nodes.  
Resulting in  $O(N^2)$  messages.

Naive Method

1. First C identities becomes Directory Servers.
2. Each node after establishing identity sends identity to nearest Directory Server.
3. Each Directory Server sends the committee/shard lists to N nodes, once each committee attains C or more nodes.





# Elastico


## Overlay Setup

How to get identities from other nodes?

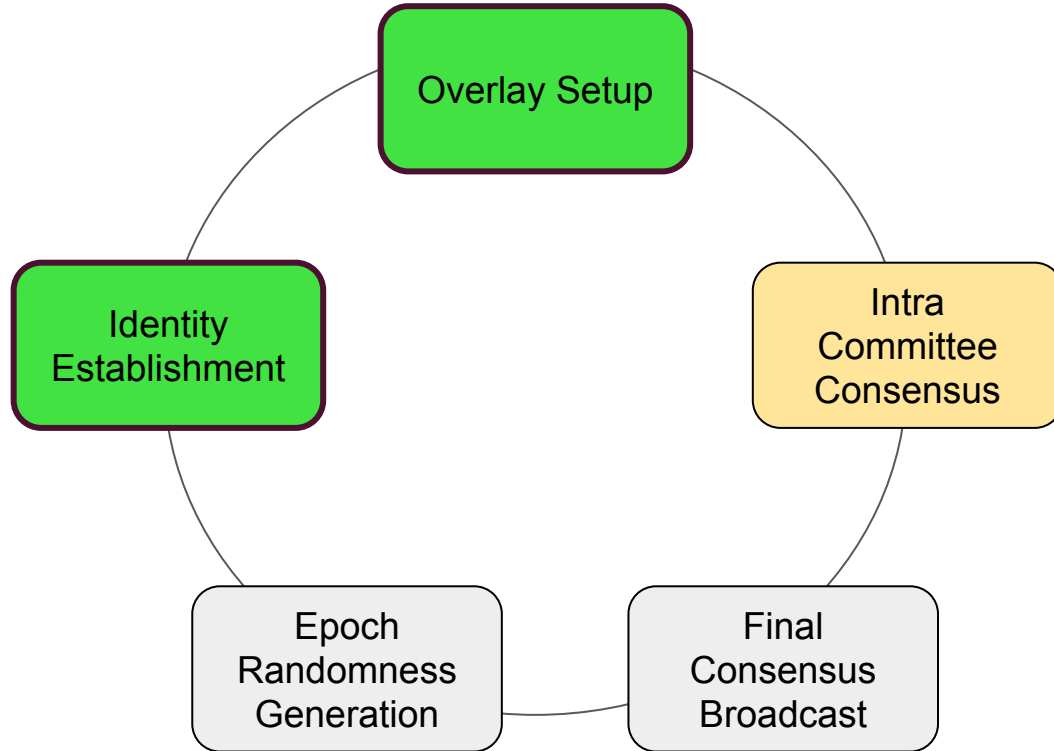
Broadcast N identity to N nodes.  
Resulting in  $O(N^2)$  messages.

Naive Method

1. First C identities becomes Directory Servers.
2. Each node after establishing identity sends identity to nearest Directory Server.
3. Each Directory Server sends the committee/shard lists to N nodes, once each committee attains C or more nodes.



# Elastico



# Elastico

## *Intra-Committee Consensus*

- **Assumption** of  $\frac{1}{3}$  malicious nodes in each committee.
- Run a **Byzantine agreement protocol** (PBFT or PolyByz).
- Number of messages in each committee =  $O(C^2)$

[C is the Committee Size]

# Elastico

## *Intra-Committee Consensus*

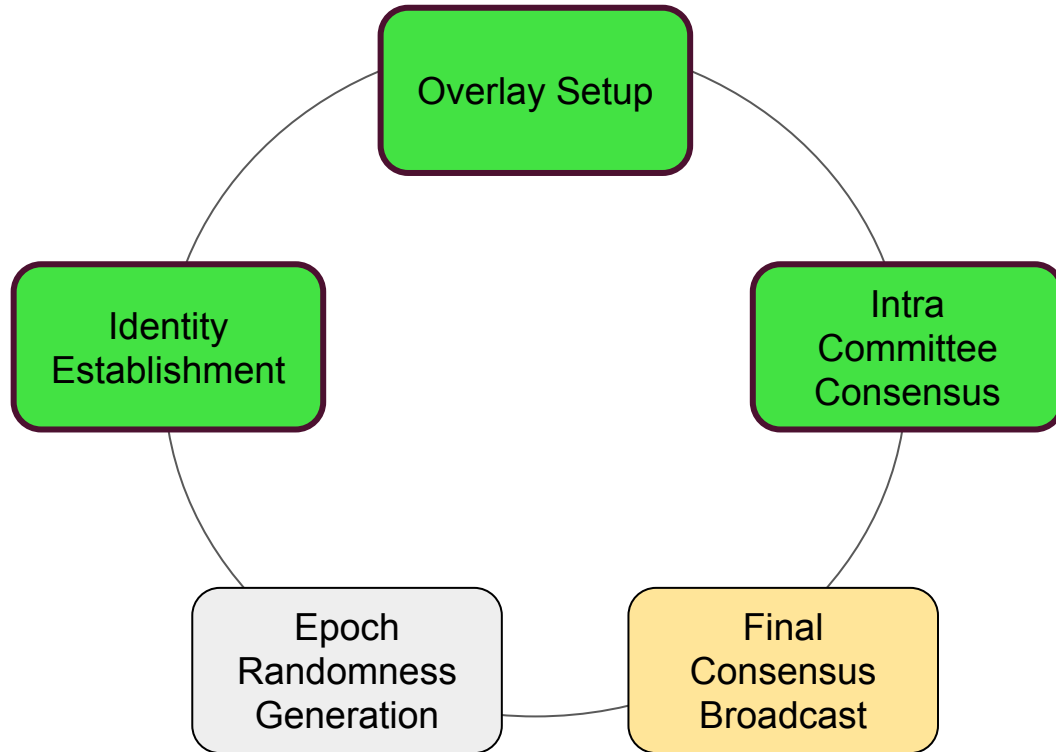
Better when  
C is more

- **Assumption** of  $\frac{1}{3}$  malicious nodes in each committee.
- Run a **Byzantine agreement protocol** (PBFT or PolyByz).
- Number of messages in each committee =  **$O(C^2)$**

**[C is the Committee Size]**

Better when  
C is less

# Elastico

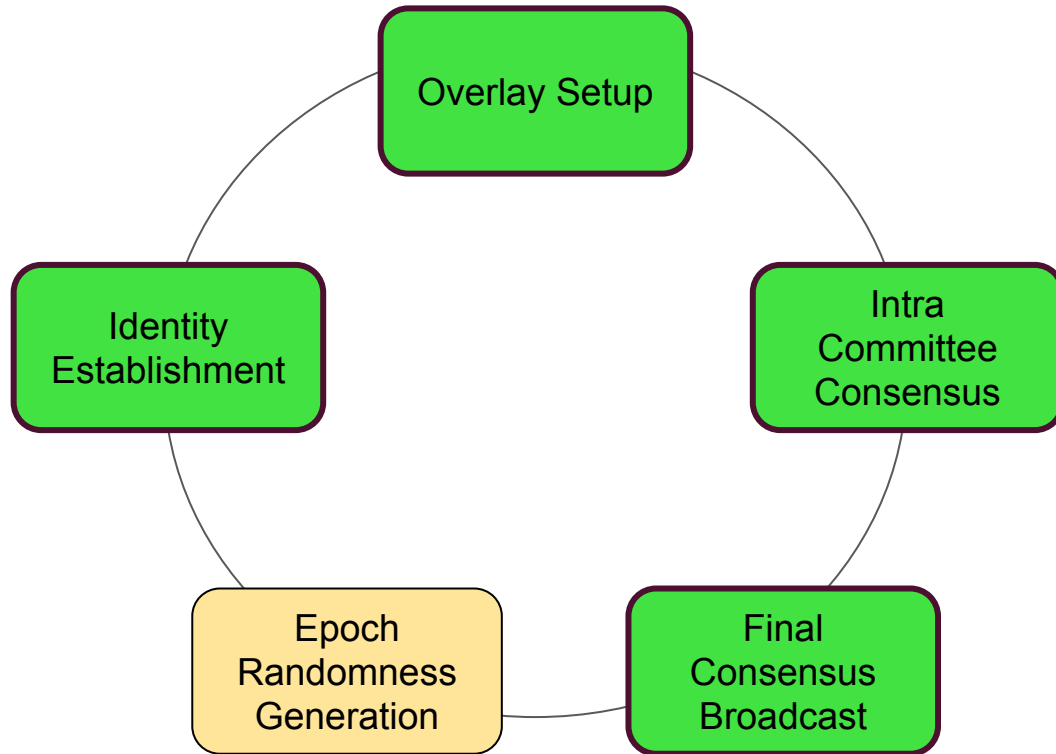


# Elastico

## *Final Consensus Broadcast*

- Form a **final committee** (consensus committee) to reach final consensus.
- The nodes constituting the final committee comprise **all members** possessing a **fixed *s-bit* committee ID**.
- **Run BFT**, and **Broadcast** the final block to everyone.

# Elastico



# Elastico

## *Epoch Randomness Generation (Commit-And-XOR)*

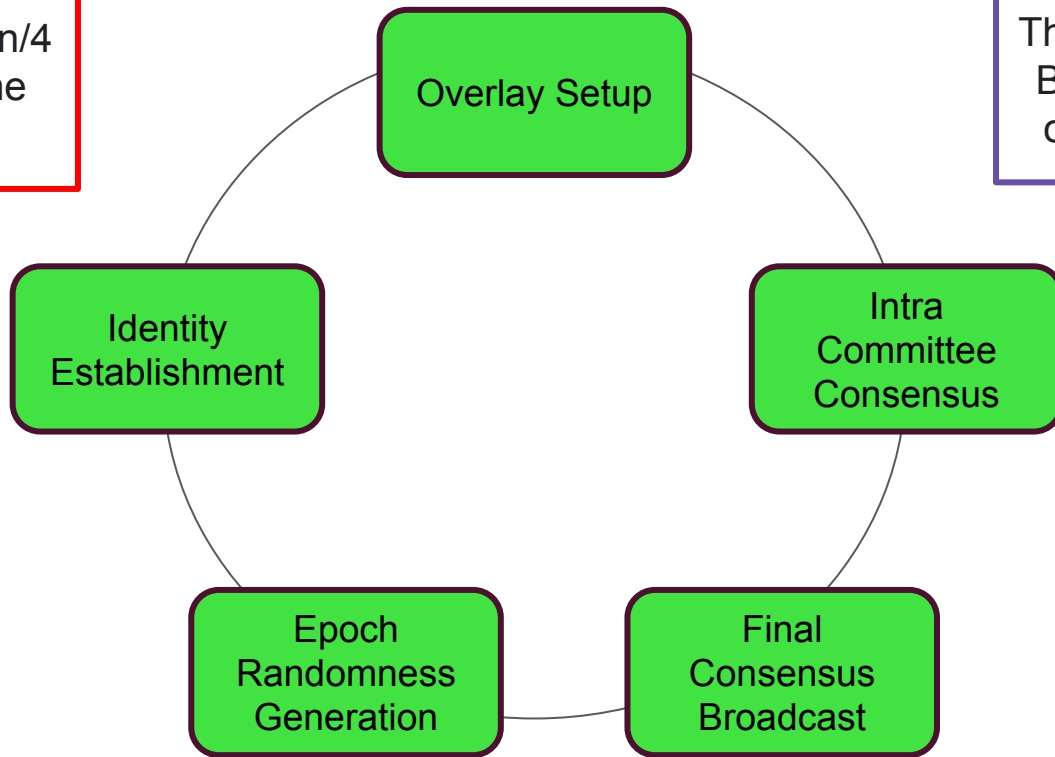
- Each member of the final committee pick a random " $R_i$ "
- Include  $H(R_i)$  in final block.
- Broadcast  $R_i$  with final block.
- Nodes XOR any set of  $2C/3$  " $R_i$ " to get randomness.



# Elastico

Tolerates up to  $f < n/4$   
adaptive byzantine  
adversaries

Throughput of 6x than  
BTC with 16 Shards  
of 100 nodes each.



# Elastico

## Limitations

- **PoW method** used for identification is susceptible to bias since miners has the ability to drop blocks, thus **influencing randomness**.
- The distributed commit-and-xor scheme is not a perfectly unbiased randomness generator.
- **High failure rate**, with total fault tolerance of 25%. Can be mitigated by increasing the consensus group but at the expense of communication overhead.
- No Cross-Shard transaction. Global state is maintained in each shard.  
**High storage overhead**.

# RapidChain [CCS '18]

*Improved Security and Throughput over Elastico*



- Tolerates up to  $f < n/3$  in total network and  $f < n/2$  in intra-shard, where  $f$  is byzantine adversaries.
- Each node exclusively stores a fraction of  $1/k$  of the ledger [ $k$  signifies the number of shards].
- Cross-Shard Protocol for efficient cross-shard transaction.
- Decentralized Bootstrapping - Does not assume the existence of an initial randomness.

# RapidChain

## *Decentralized Bootstrapping*

RapidChain proposes a decentralized bootstrapping in the form of **sampler-graph election network**, with only a **hardcoded seed and some network settings**.

In such an election network, participating **validators are uniformly distributed** into a **few groups**, within each of which a **PoW-based result is computed by each member** based on the **randomness generated by the VSS-based DRG protocol** and its identification ID.

Based on the result, a subgroup can be obtained for each group. Finally, a unique root group (it **randomly selects the members of the reference committee**) can be obtained with 50% honest majority (high probability), when this process is iterated. Consequently, the communication overhead can be improved from  $\Theta(n^2)$  to  $O(n\sqrt{n})$  with  $n$  denoting the total number of participating validators.

# RapidChain

## *Epoch Reconfiguration*

RapidChain implements a VSS-based distributed random generation (DRG) protocol to agree on an unbiased randomness.

The proposed DRG protocol by RapidChain, in fact, only implements a basic VSS-shares scheme, where all participating validators can reconstruct the final randomness  $r$  by the share of  $r$  (the share equals to  $\sum_{l=1}^m \rho_{lj}$  calculated by other validators except validator- $j$ ) received from other validators. Note that,  $\rho \in \mathbb{F}_p$  denoting a finite field of prime order  $p$ , and  $m$  denotes the size of the reference committee. As a result, the DRG protocol encounters a similar issue to that of any other typical VSS scheme, i.e., non-scalable (even though it suits with the 50% BFT in small-sized shards).

# RapidChain

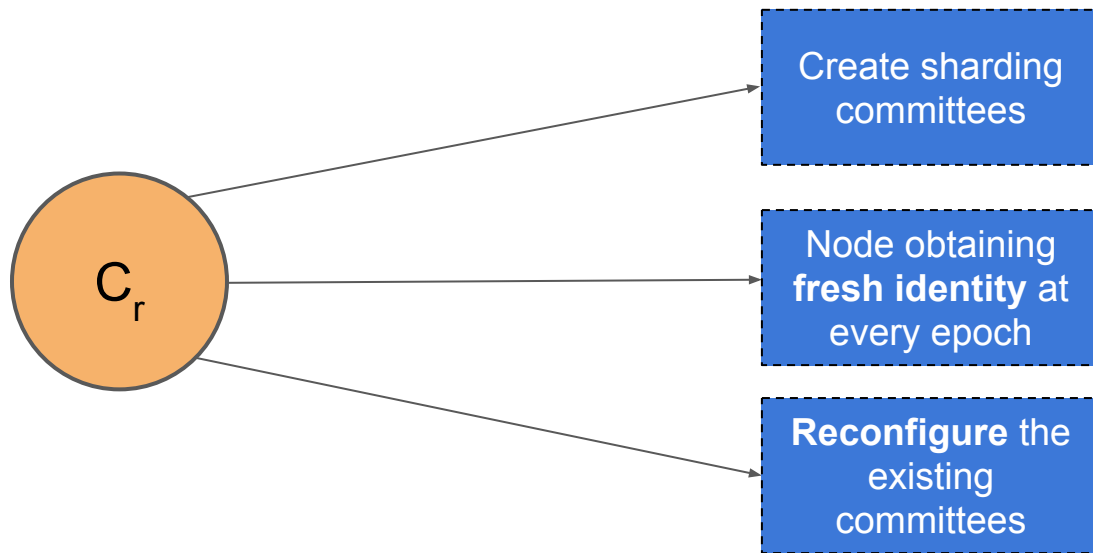
## *Intra-Consensus*

- Runs an autonomous **pre-scheduled scheme** within a shard to **agree on a timeout  $\Delta$** .
- Based on the timeout, the consensus speed can be adjusted by the system to **prevent the asynchronization**.
- A non-responsive synchronous (with constant rounds) BFT-based consensus protocol with **FT of 50%** can be used.

# RapidChain

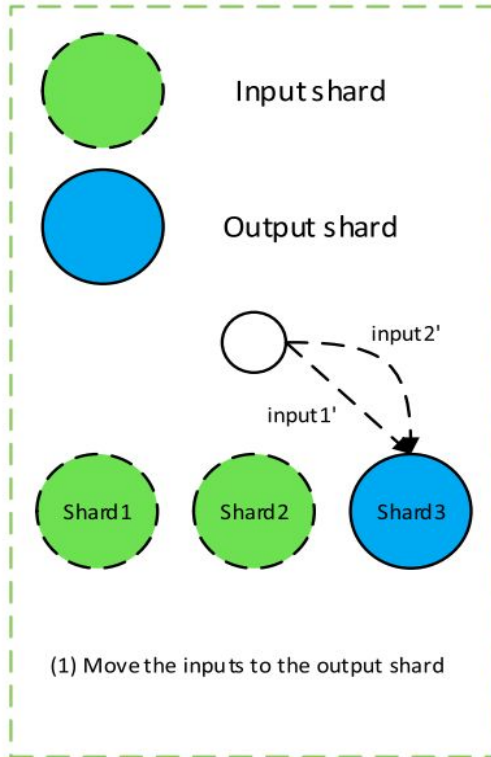
## *Reference Committee ( $C_r$ )*

At the conclusion of each epoch 'i',  
 $C_r$  generates a fresh randomness denoted as  $r_{i+1}$ ,  
serving as the epoch randomness for the subsequent epoch  $i + 1$ .



# RapidChain

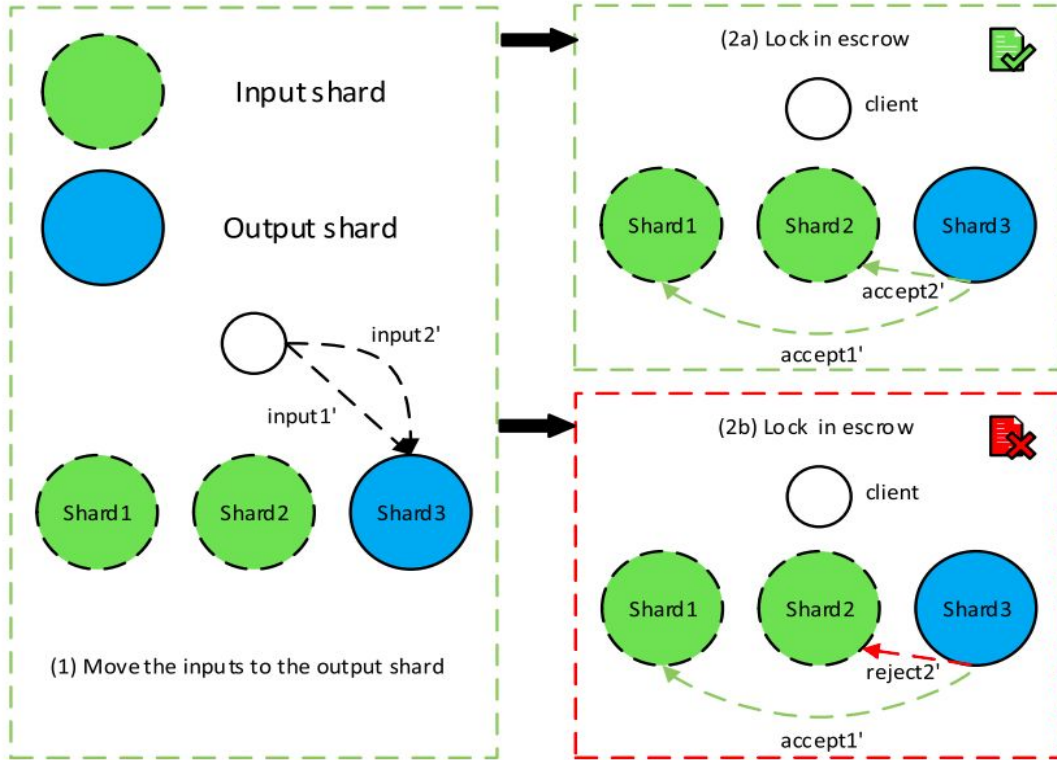
## *Cross-Shard Transaction*





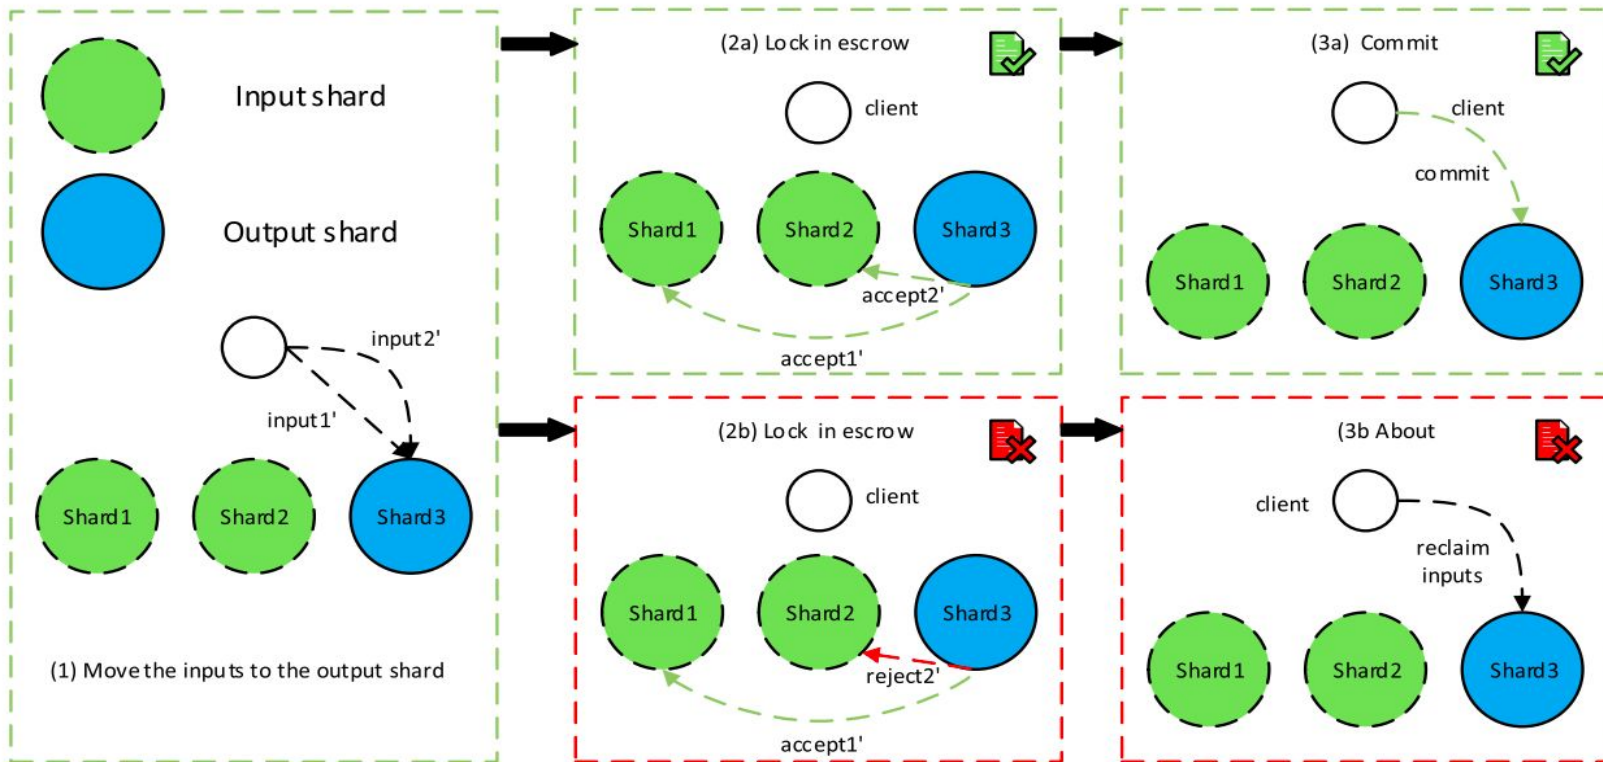
# RapidChain

## *Cross-Shard Transaction*




# RapidChain

## Cross-Shard Transaction



# RapidChain

## *Cross-Shard Transaction*



$$\text{"TX"} = \langle (I_1, I_2), O \rangle$$


Input  $I_1$  (from shard  $S_1$ )

Input  $I_2$  (from shard  $S_2$ )

Output  $O$  (to shard  $S_3$ ).

# RapidChain

## *Cross-Shard Transaction*



$S_3$ , creates 3 New Transactions:

$$TX_1 = \langle I_1, I_1' \rangle$$



$$TX_2 = \langle I_2, I_2' \rangle$$

$$TX_3 = \langle (I_1', I_2'), O \rangle$$

# RapidChain

## *Limitations*

→ Rapidchain's **inability** to maintain **isolation** is evident in certain scenarios.



Only one of these  
two succeeds

If  $TX_{xyz}$  fails while  $TX_2$  succeeds but  $TX_1$  fails, none of the transactions would be successfully processed.

# RapidChain

## *Throughput Comparison with Elastico and Omniledger*

Protocol	# Nodes	Resiliency	Complexity <sup>1</sup>	Throughput	Latency	Storage <sup>1</sup>	Shard Size	Time to Fail
Elastico [47]	$n = 1,600$	$t < n/4$	$\Omega(m^2/b+n)$	40 tx/sec	800 sec	1x	$m = 100$	1 hour
OmniLedger [42]	$n = 1,800$	$t < n/4$	$\Omega(m^2/b+n)$	500 tx/sec	14 sec	1/3x	$m = 600$	230 years
OmniLedger [42]	$n = 1,800$	$t < n/4$	$\Omega(m^2/b+n)$	3,500 tx/sec	63 sec	1/3x	$m = 600$	230 years
RapidChain	$n = 1,800$	$t < n/3$	$O(m^2/b+m \log n)$	4,220 tx/sec	8.5 sec	1/9x	$m = 200$	1,950 years
RapidChain	$n = 4,000$	$t < n/3$	$O(m^2/b+m \log n)$	<b>7,380 tx/sec</b>	<b>8.7 sec</b>	<b>1/16x</b>	<b><math>m = 250</math></b>	<b>4,580 years</b>

Zamani, Mahdi, Mahnush Movahedi, and Mariana Raykova. "Rapidchain: Scaling blockchain via full sharding."

# RapidChain


## *Limitations*

- With the increment of shards, nearly **all transactions evolve into cross-shard transactions**. Batch processing of cross-shard transactions provides partial alleviation.
- The **latency increases by 1.5 times if 10 new nodes join** the network for each shard.
- There is **no provision for adaptive-based sharding** in RapidChain. Restarting the entire network is imperative to modify the number of shards.
- RapidChain's methodology proves **inadequate for non-UTXO distributed transactions** as it violates both atomicity and isolation principles.

# Monoxide [NSDI'19]

Addresses Hash Power dilution by a novel PoW based mining technique

- Adopts the concept of asynchronous consensus zones, where each zone represents a shard.
- Follows **Account based transaction model**.
- A novel **Chu-Ko-Nu mining technique** rooted in the Nakamoto-based Proof-of-Work. (Fault Tolerance of 50%)
- Leverages **Relay Transactions** for cross-shard communication.





# Monoxide

## Chu-Ko-Nu Mining

- Miners will be able to mine on multiple chains of different shards at the same time by solving one PoW target.




.....



# Monoxide

## Chu-Ko-Nu Mining

- Miners will be able to mine on multiple chains of different shards at the same time by solving one PoW target.



# Monoxide

## Chu-Ko-Nu Mining

- The **Merkle Patricia tree** (MPT) serves as the storage structure for all proposed blocks from multiple shards.
- $MPT_M$  denotes the MPT root consisting of **all proposed blocks of each involved shard**, i.e.,  $[B_0, B_1, \dots, B_{n-1}]$  if  $k = n$ .

The consensus algorithm in Monoxide is expressed as follows:

$$H(\phi || H(x || MPT_M)) \leq \epsilon$$

H = Hash Function

$\phi$  = Random Nonce

x = Block header content

$\epsilon$  = PoW target difficulty.

# Monoxide

## Chu-Ko-Nu Mining

- The Chu-ko-nu mining algorithm enables miners in any shard to generate multiple blocks upon successfully solving the hash puzzle, limited to one block per shard.

This amplifies the computing power of miners in each shard, thereby requiring an equivalent level of computing power to attack a single shard as attacking the entire system.

# Monoxide

## Chu-Ko-Nu Mining

- Monoxide mandates that most miners in each shard engage in Chu-ko-nu mining (i.e.,  $k \rightarrow n$ ) to satisfy  $P^*(k/n) \approx P$ .

“P” denotes the total amount of mining power,  
“n” signifies the number of shards.

- This requirement is essential because scattered mining power may not ensure the security of a single shard.

# Monoxide

## Cross-Zone Transactions

Transfer  $x$  tokens from user **A** to user **B** in different zones

$$A \leftarrow A - x, (A \geq x)$$

Conditional Operation  
Order-dependent



Execute in Zone A  
Update A's balance

$$B \leftarrow B + x$$

Unconditional Operation  
Order-independent

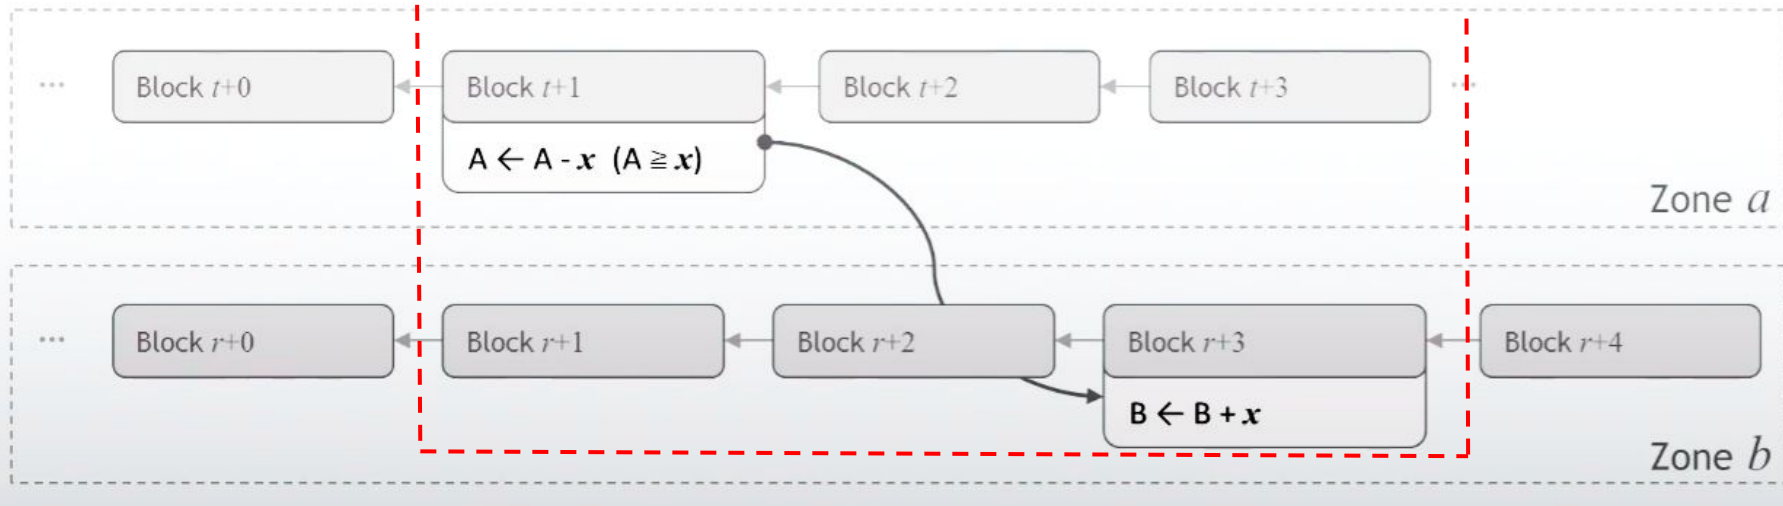


Execute in Zone B  
Update B's balance



# Monoxide

## Cross-Zone Transactions



# Monoxide

## Limitations

- If miners only mine on “k” out of “n” shards, where  $(k \ll n)$ 
  - ◆ The factor expected to amplify the effective mining power will be too small to secure the mining process, hence reducing the attack cost.
- Rational miners tend to mine on all “n” shards to reap the maximum profit
  - ◆ May also result in the power centralization due to the huge cost of bandwidth, disk storage, and computing processors that only the professional mining facilities can afford.



# GearBox [CCS' 22]

Adaptively Change Shard Size based on  
Adversary Population

- Segregates the concepts of **liveness and safety** within shard consensus.



# GearBox [CCS' 22]

Adaptively Change Shard Size based on Adversary Population

- Segregates the concepts of **liveness and safety** within shard consensus.
- **Dynamic adjustment of shard parameters** to achieve optimal efficiency given the current corruption ratio of the system.



# GearBox [CCS' 22]


Adaptively Change Shard Size based on Adversary Population

- Segregates the concepts of **liveness and safety** within shard consensus.
- **Dynamic adjustment of shard parameters** to achieve optimal efficiency given the current corruption ratio of the system.
- Shards identified as not meeting liveness requirements are subject to **resampling with increased shard size** and liveness tolerance.



# GearBox

## Chernoff Bound



- Sampling a random subset may not always provide the expected corruption level.
- Chernoff bound offers a means to limit this additional corruption ( $\epsilon$ ) to an arbitrarily small value when the random subset is sufficiently large.

# GearBox

## Safety-Liveness Dichotomy

### Safety Threshold (S)

How much adversary can my network handle without sacrificing safety?

### Liveness Threshold (L)

How much adversary can my network handle without sacrificing liveness?



# GearBox

Safety-Liveness Dichotomy

**When  $S = L$**

# GearBox

## Safety-Liveness Dichotomy

### When $S = L$

#### Synchronous Model

$$S < 1 - L$$

$$S + L < 1$$

$$L, S < \frac{1}{2}$$

#### Partially Synchronous Model

$$S < 1 - 2L$$

$$S + 2L < 1$$

$$L, S < \frac{1}{3}$$

# GearBox

## Safety-Liveness Dichotomy

### When $S = L$

#### Synchronous Model

$$\begin{aligned} S + L &< 1 \\ L, S &< \frac{1}{2} \end{aligned}$$

#### Partially Synchronous Model

$$\begin{aligned} S + 2L &< 1 \\ L, S &< \frac{1}{3} \end{aligned}$$


To ensure liveness during block commitment, the block must be endorsed by a fraction of nodes equal to  $1 - L$ , where the proportion of corrupted nodes is no more than  $L$ .



# GearBox

## Safety-Liveness Dichotomy


In the case of **two conflicting blocks** -



If “ $1-L$ ” of the nodes have endorsed both blocks, there will be an intersection among the endorsers indicating that **more than “ $S$ ” parties have endorsed both blocks**.

# GearBox

## Safety-Liveness Dichotomy



As “S” represents the maximum number of malicious parties, **there must be at least one honest node that endorses both blocks**, thus leading to the detection of the conflict.

# GearBox

## Safety-Liveness Dichotomy

To run a consensus based on the dichotomy of this partially synchronous network  $S + 2L < 1$

*The protocol needs to continuously change in  $L$  and  $S$  based on the malicious nodes present in the network.*

# GearBox

Gears for Changing Shard Size

Presents 6 gears to change the configuration of the network based on the situation.

# GearBox

## Gears for Changing Shard Size

Presents 6 gears to change the configuration of the network based on the situation.

Gear	1	2	3	4	5	6
S	39%	49%	59%	69%	79%	89%
L	30%	25%	20%	15%	10%	5%
Shard size to ensure safety	1713	462	207	116	75	51

David, Bernardo, et al. "Gearbox: Optimal-size shard committees by leveraging the safety-liveness dichotomy."

# GearBox

## Gears for Changing Shard Size

Presents 6 gears to change the configuration of the network based on the situation.

Gear	1	2	3	4	5	6
S	39%	49%	59%	69%	79%	89%
L	30%	25%	20%	15%	10%	5%
Shard size to ensure safety	1713	462	207	116	75	51

David, Bernardo, et al. "Gearbox: Optimal-size shard committees by leveraging the safety-liveness dichotomy."

Will this get into Liveness Attack?

# GearBox


Shard Liveness Monitoring

Challenge: Monitor the liveness of each shard

# GearBox

## Shard Liveness Monitoring

Challenge: Monitor the liveness of each shard





# GearBox

## Limitations

### **Issues with overlapping shards:**


Since the protocol adjusts shard sizes without altering the total number of shards. These overlapping shards could lead to duplicated workloads, requiring additional transactions involving multiple shards. Consequently, there might be a rise in cross-shard transactions and a decrease in parallelism.

# GearBox

## Limitations

### Issues with overlapping shards:

Since the protocol adjusts shard sizes without altering the total number of shards. These overlapping shards could lead to duplicated workloads, requiring additional transactions involving multiple shards. Consequently, there might be a rise in cross-shard transactions and a decrease in parallelism.



# GearBox

## Limitations

Simply altering shard sizes in both directions to accommodate dynamic changes in  $L$  is not feasible.

# GearBox

## Limitations

Simply altering shard sizes in both directions to accommodate dynamic changes in  $L$  is not feasible.

Attackers could **exploit this method to trigger a loop of "switching gears"**, resulting in frequent and expensive shard adjustments.

# Reticulum [NSDI' 2024]

A two-phase framework that dynamically adjusts transaction throughput in response to real-time adversarial attack.



# Reticulum [NSDI' 2024]

A two-phase framework that dynamically adjusts transaction throughput in response to real-time adversarial attack.



- It comprises 'control' and 'process' shards organized into two layers corresponding to the two phases.

# Reticulum [NSDI' 2024]

A two-phase framework that dynamically adjusts transaction throughput in response to real-time adversarial attack.



- It comprises 'control' and 'process' shards organized into two layers corresponding to the two phases.
- Process shards are subsets of control shards.

# Reticulum [NSDI' 2024]

A two-phase framework that dynamically adjusts transaction throughput in response to real-time adversarial attack.




- It comprises 'control' and 'process' shards organized into two layers corresponding to the two phases.
- Process shards are subsets of control shards.
- Process shards leverages Safety while Control Shards leverages Liveness.



# Reticulum


## Two-Phase-Shard Consensus



# Reticulum

## Two-Phase-Shard Consensus


→ Each **Process Shard** expected to **include at least one honest node** with high confidence.



# Reticulum

## Two-Phase-Shard Consensus


- Each **Process Shard** expected to include at least one honest node with high confidence.
- Process Shards follows **unanimous voting** to engage fewer nodes in the First Phase.



# Reticulum

## Two-Phase-Shard Consensus


- Each **Process Shard** expected to include at least one honest node with high confidence.
- Process Shards follows *unanimous voting* to engage fewer nodes in the First Phase.
- Each **Control Shard** expected to consist of a majority of honest nodes with high confidence.



# Reticulum

## Two-Phase-Shard Consensus

- Each **Process Shard** expected to **include at least one honest node** with high confidence.
- **Process Shards** follows **unanimous voting** to engage fewer nodes in the First Phase.
- Each **Control Shard** expected to **consist of a majority of honest nodes** with high confidence.
- Control Shard then **finalizes the decision** made in the first phase.



# Reticulum

## The First Phase

- **Unanimous voting** is required within each process shard to determine a block validity.

# Reticulum


## The First Phase

- **Unanimous voting** is required within each process shard to determine a block validity.
- If even one adversarial node is present within the process shard, consensus cannot be reached.

# Reticulum

## The First Phase

- **Unanimous voting** is required within each process shard to determine a block validity.
- If even one adversarial node is present within the process shard, consensus cannot be reached.






# Reticulum

## The First Phase

**Process Shard fails even if one malicious node is there.**


- **Unanimous voting** is required within each process shard to determine a block validity.
- If even one adversarial node is present within the process shard, consensus cannot be reached.
- Hence,  $L = 0$ , while  $S < 100\%$ .



# Reticulum

## The First Phase

**Process Shard fails even if one malicious node is there.**



- **Unanimous voting** is required within each process shard to determine a block validity.
- If even one adversarial node is present within the process shard, consensus cannot be reached.
- Hence,  $L = 0$ , while  $S < 100\%$ .
- The first phase must be concluded within a predetermined time limit denoted as  $T_1$ .

# Reticulum

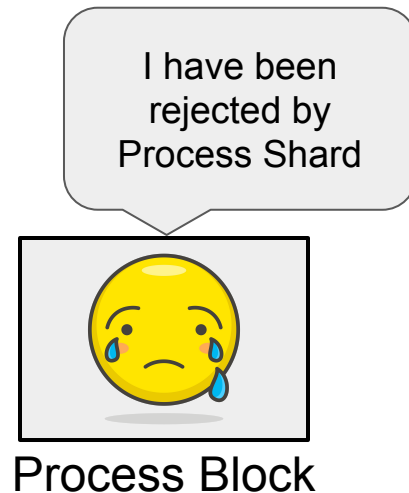
## The Second Phase

- “Safety Net” for blocks that did not receive unanimous approval.

# Reticulum

## The Second Phase

→ “Safety Net” for blocks that did not receive unanimous approval.




# Reticulum

## The Second Phase

→ “Safety Net” for blocks that did not receive unanimous approval.

I have been  
rejected by  
Process Shard



Process Block



Don't Worry!  
I am here.


Control Shard

# Reticulum

## The Second Phase

- “Safety Net” for blocks that did not receive unanimous approval.
- Leader from the Control Shard proposes a Control Block with details about the Process Blocks from each shard under its jurisdiction.

I have been rejected by Process Shard



Process Block



Don't Worry!  
I am here.


Control Shard

# Reticulum

## The Second Phase

- “**Safety Net**” for blocks that did not receive unanimous approval.
- **Leader** from the Control Shard **proposes a Control Block** with details about the Process Blocks from each shard under its jurisdiction.
- The **second phase must be completed within T2**, which, unlike T1, is dynamically adjusted based on the number of successful process shards within the control shard.

I have been  
rejected by  
Process Shard



Process Block



Don't Worry!  
I am here.

Control Shard

# Reticulum

## Drawbacks

- Uses a **static sharding scheme**, i.e. no new members can be added in runtime and no respawn of shards can take place. A **system reboot is necessary for fresh node addition**.
- **No proof for cross-shard isolation is given in the paper**. It could fail the same way RapidChain's isolation failed.



# Comparison Table

		Elastico	Omniledger	RapidChain		Monoxide	Ethereum 2.0	GearBox	Reticulum
Network Model		Partial Sync	Partial Sync	Intra	Sync	Partial Sync	Partial Sync	Async	Sync
				Total	Partial Sync				
Threat Model		Arbitrary Behaviour of Attacker, Round-Adaptive Adversary	Arbitrary Behaviour of Attacker, Mildly-Adaptive Adversary	Arbitrary Behaviour of Attacker, Slowly-Adaptive Adverasary		Similar to Bitcoin	Arbitrary Behaviour of Attacker, Uncoordinated Majority	Static Adaversary	Arbitrary Behaviour of Attacker, Adaptive but Upper-Bounded Adversary
Fault Tolerance	Intra	33%	33%	50%		50%	33%	Adaptive	50%
	Total	25%	25%	33%		50%	33%	33%	33%
Intra-Consensus Protocol		PBFT	ByzCoinX	50% BFT		PoW-based (Chu-Ko-Nu) Mining	BFT-Based PoS	BFT with varied FTR	BFT with Unanimous Majority
Transaction Structure		UTXO	UTXO	UTXO		Account	Account	UTXO / Account	N/A
Cross-Shard Communication		N/A	Atomix Involves User (Lock / Unlock)	Divides into 3 new Txns (Lock / Unlock)		Relay Transaction (Lock Free)	Receipt Based (Lock / Unlock)	Customized Atomix (Lock / Unlock)	Customized Rapidchain Protocol (Lock / Unlock)
Additional Security Chain		Global Ledger	Global Identity Blockchain	Reference Committee		N/A	Beacon Chain	Control Chain	Committee Shard, Public Communication Chain
Throughput		13.5 times of Bitcoin's TPR, 1600 Nodes	4000 TPR, 1800 Nodes, 25% FTR, Shard Size=600	7300 TPR, 4000 Nodes, 33% FTR, Shard Size=250		1000 times of Bitcoin's TPR, 48000 Nodes	100000 TPR on mainnet	32000 TPR, Shard Size=82. 1561 TPR, Shard Size=2264	2000 TPR with 50% Process Shard Acceptance. 7000 TPR with 95% Process Shard Acceptance.
Adaptive Shard Size		No	No	No		No	No	Yes	No

# Issues and Future Works

Limiting  
Cross-Shard  
Transactions

Reducing  
Transaction  
Latency

Pruning  
Shard's  
Ledger

Data Validity  
Problem

Data  
Availability  
Problem

# References

- [1] Luu, Loi, et al. "A secure sharding protocol for open blockchains." Proceedings of the 2016 ACM SIGSAC conference on computer and communications security. 2016.
- [2] Kokoris-Kogias, Eleftherios, et al. "Omniledger: A secure, scale-out, decentralized ledger via sharding." 2018 IEEE symposium on security and privacy (SP). IEEE, 2018.
- [3] Zamani, Mahdi, Mahnush Movahedi, and Mariana Raykova. "Rapidchain: Scaling blockchain via full sharding." Proceedings of the 2018 ACM SIGSAC conference on computer and communications security. 2018.
- [4] Wang, Jiaping, and Hao Wang. "Monoxide: Scale out blockchains with asynchronous consensus zones." 16th USENIX symposium on networked systems design and implementation (NSDI 19). 2019.
- [5] David, Bernardo, et al. "Gearbox: Optimal-size shard committees by leveraging the safety-liveness dichotomy." Proceedings of the 2022 ACM SIGSAC Conference on Computer and Communications Security. 2022.
- [6] Xu, Yibin, et al. "A Two-Layer Blockchain Sharding Protocol Leveraging Safety and Liveness for Enhanced Performance." arXiv preprint arXiv:2310.11373 (2023).

# Backup Slides

# Omniledger [SP'2018]



*Improves on the generation of randomness and proposes Atomix Protocol for Cross-Shard Transaction*

- ❖ **Better epoch randomness** - Combination of RandHound and Algorand-based Verifiable Random Function (VRF)
- ❖ **Better consensus algorithm** - ByzCoinX
- ❖ Atomix Protocol for **efficient cross shard** transaction.
- ❖ **Tolerates up to  $f < n/4$**  adaptive byzantine adversaries.

# Omniledger

## *Epoch Randomness and Shard Validator Assignment*

- **Challenge:** Unbiasable, unpredictable and scalable shard validator assignment

# Omniledger

## *Epoch Randomness and Shard Validator Assignment*


- **Challenge:** Unbiasable, unpredictable and scalable shard validator assignment
- **Solution:** Combine VRF-based lottery and unbiasable randomness protocol for sharding

# Omniledger


## *Shard Validator Assignment*

- **Challenge:** Unbiasable, unpredictable and scalable shard validator assignment
- **Solution:** Combine VRF-based lottery and unbiased randomness protocol for sharding


1. Temp. leader election  
via VRFs (biasable)



2. Randomness generation  
via RandHound\* (unbiasable)



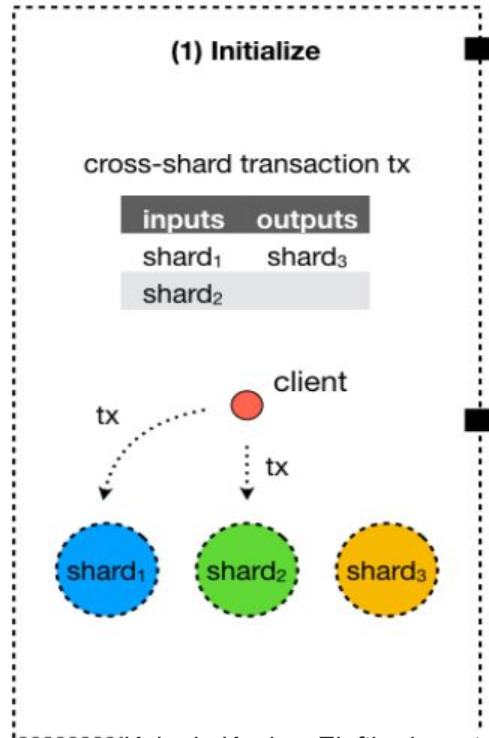
3. Shard assignment  
(using  $rnd_e$ )





# Omniledger

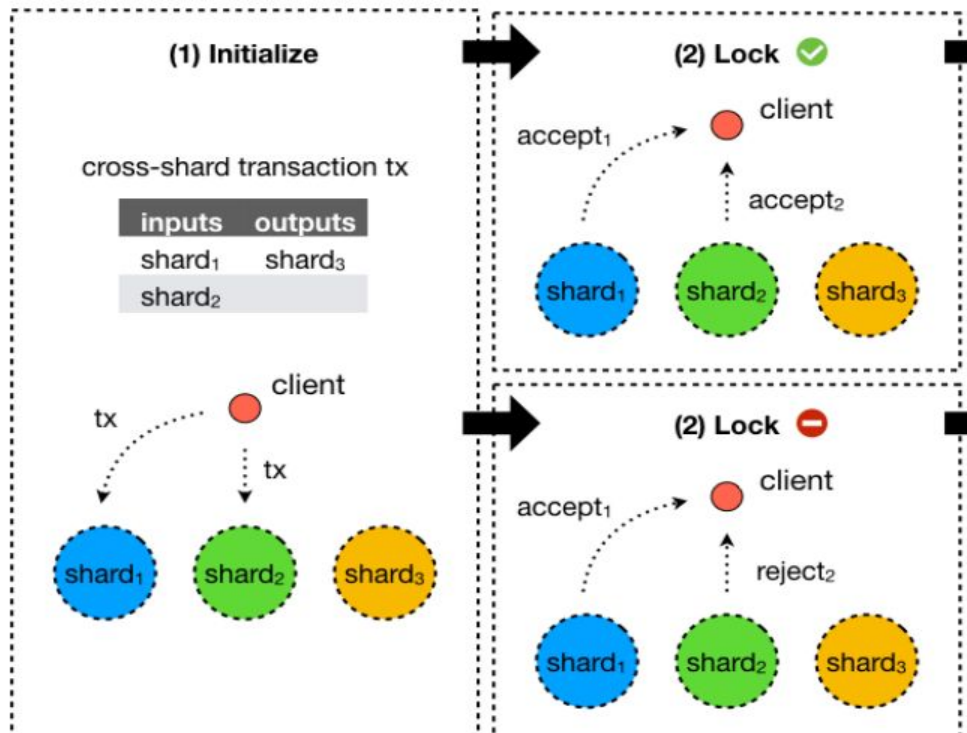
## *Atomix - Cross Shard Transaction*



Kokoris-Kogias, Eleftherios, et al. "Omniledger: A secure, scale-out, decentralized ledger via sharding."

# Omniledger

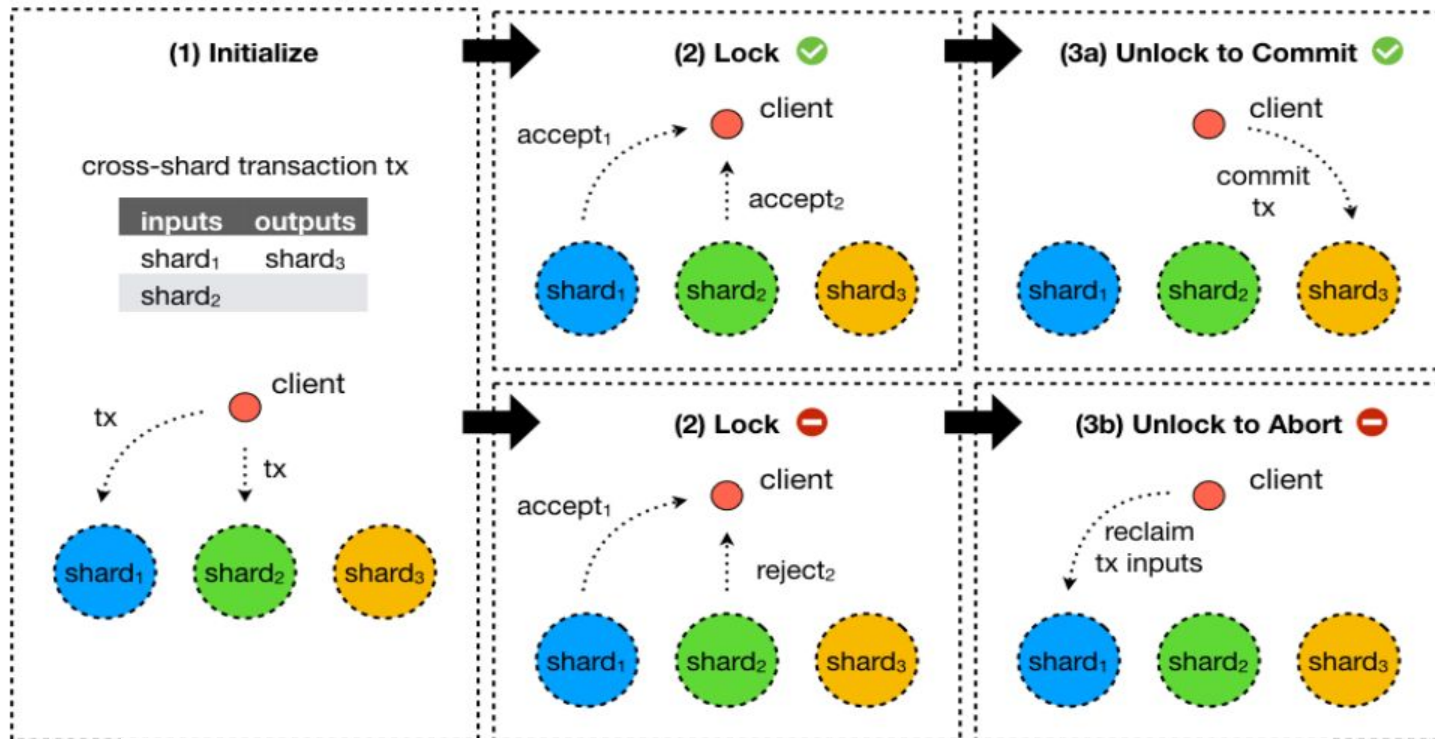
## *Atomix - Cross Shard Transaction*



Kokoris-Kogias, Eleftherios, et al. "Omniledger: A secure, scale-out, decentralized ledger via sharding."

# Omniledger

## *Atomix - Cross Shard Transaction*




# Omniledger

## *Limitations*

- The combination of RandHound and VRF suffers from the **reliance on a third-party initial randomness** pre-defined in the genesis block.
- Atomix Protocol **requires a client to actively participate** in cross-shard transactions.
- This client-driven protocol suffers from **indefinite blocking if the client is malicious**. If Payee is a client that coordinates a transaction that transfers funds from a payer's account. A malicious payee may pretend to crash indefinitely during the lock/unlock protocol, hence, the payer's funds are locked forever.

# Ethereum 2.0



Sharding is splitting up the Ethereum blockchain so that subsets of validators are only responsible for a fraction of the total data. This was originally intended to be the way for Ethereum to scale. However, layer 2 rollups have developed much faster than expected and have provided a lot of scaling already, and will provide much more after Proto-Danksharding is implemented. This means "shard chains" are no longer needed and have been dropped from the roadmap.

# Ethereum 2.0

## Dank-Sharding

- Rollups are a way to scale Ethereum by batching transactions off-chain and then posting the results to Ethereum.
- A Rollup is essentially composed of two parts: Data and Execution check.
- The data is the full sequence of transactions that is being processed by a rollup.
- The execution check is the re-execution of those transactions by some honest actor (a "prover") to ensure that the proposed state change is correct

# Ethereum 2.0

## Dank-Sharding

- To perform the execution check, the transaction data has to be available for long enough for anyone to download and check.
- Rollups uses **CALLDATA** for posting transactions, which proves costly as it is processed by all Ethereum nodes and permanently stored on-chain.
- Danksharding introduces **Data Blobs**. Rollups post commitments to their transaction data on-chain and also make the actual data available in these data blobs.

# Ethereum 2.0

## Dank-Sharding

- The data in these blobs is not accessible to the EVM and is **automatically deleted after a fixed time period**.
- At the node-level, the **blobs of data are held in the consensus client**. The consensus clients attest that they have seen the data and that it has been propagated around the network.
- Validators will **have enough time to check the commitments and challenge data of a rollup transaction** while the consensus client can provide proof for it.